

OINK Users Manual

OINK and MapReduce–MPI Library

<http://www.sandia.gov/~sjplimp/mapreduce.html> – Sandia National Laboratories

Copyright (2009) Sandia Corporation. This software and manual is distributed under the BSD License.

Table of Contents

OINK Documentation.....	1
Goals of OINK.....	1
Contents of OINK Manual.....	2

OINK Documentation

Version info:

The OINK "version" is the date when it was released, such as 1 Feb 2011. OINK and MR-MPI library are updated continuously. Whenever we fix a bug or add a feature, we release it immediately, and post a notice on [this page of the WWW site](#). Each dated copy of OINK contains all the features and bug-fixes up to and including that version date. The version date is printed to the screen and log file every time you run OINK. It is also in the file oink/version.h and in the MR-MPI directory name created when you unpack a tarball.

- If you browse the HTML doc pages for OINK on the MR-MPI WWW site, they always describe the most current version of OINK.
- If you browse the HTML doc pages for OINK included in your tarball, they describe the version you have.
- The [PDF file](#) for OINK on the WWW site or in the tarball is updated about once per month. This is because we don't want it to be part of every patch.

OINK is a simple scripting wrapper around the "MapReduce-MPI [library](#)", and also provides a easy-to-use development framework for writing new MapReduce algorithms and codes. Like the MR-MPI library, OINK will run on any platform (serial or parallel) that supports [MPI](#). Note the MR-MPI library has its own [manual and doc pages](#).

The name OINK is meant to evoke the aroma of the [Apache Pig](#) platform which wraps the [Hadoop](#) MapReduce capabilities with its high-level Pig Latin language. Since OINK has only a tiny fraction of Pig's capability, it is more the sound of a pig, than a pig itself.

Source code for OINK and the MR-MPI library were developed at Sandia National Laboratories, a US Department of Energy facility. They are freely available for download from the [MR-MPI web site](#) and are distributed under the terms of the modified [Berkeley Software Distribution \(BSD\) License](#). This basically means they can be used by anyone for any purpose. See the LICENSE file provided with the distribution for more details.

The authors of OINK and the MR-MPI library are [Steve Plimpton](#) at and [Karen Devine](#) who can be contacted via email: sjplimp,kddevin@sandia.gov.

Goals of OINK

- (1) To allow MapReduce algorithms which call the MR-MPI library to be written with a minimum of extraneous code, to work with input/output in various forms, and to be chained together and driven via a simple, yet versatile scripting language.
- (2) To create an archive of map() and reduce() functions for re-use by different algorithms.
- (3) To provide a scripted interface to the lo-level MR-MPI library calls that can speed development/debugging of new algorithms before coding them up in C++ or another language.

We think the first two goals are largely met. See the section on [Adding Commands to OINK](#) and the [named command](#), [input](#), and [output](#) doc pages for details of the first goal. See the section on [Adding Functions to OINK](#) for details of the second goal.

The third goal, however, is only partially met. See the [MR-MPI library commands](#) doc page for its current status. The sticking point here is that in a real programming language you can pass a pointer to an arbitrary data structure to your map() or reduce() functions, but it is hard to do that from a scripting language using text input without re-inventing something like [Python](#).

Contents of OINK Manual

OINK aims to be a simple scripting interface and development environment and the lightweight documentation reflects that.

Once you are familiar with OINK, you may want to bookmark [this page](#) at `Section_scripts.html#comm` since it gives quick access to documentation for all OINK commands.

[PDF file](#) of the entire manual, generated by [htmldoc](#)

- [Building OINK](#)
 - 1.1 [Making OINK](#)
 - 1.2 [Building OINK as a library](#)
 - 1.3 [Running OINK](#)
 - 1.4 [Command-line options](#)
- [OINK Scripts](#)
 - 2.1 [Input script operation](#)
 - 2.2 [Parsing rules](#)
 - 2.3 [Input script commands](#)
- [Adding Functions to OINK](#)
 - 3.1 [Map\(\) functions](#)
 - 3.2 [Reduce\(\) functions](#)
 - 3.3 [Compare\(\) functions](#)
 - 3.4 [Hash\(\) functions](#)
 - 3.5 [Scan\(\) functions](#)
- [Adding Commands to OINK](#)
 - 4.1 [Source files for the new class](#)
 - 4.2 [Methods in the new class](#)
 - 4.3 [Calls to the OINK object manager](#)
 - 4.4 [Calling back to map\(\) and reduce\(\) functions](#)
- [Errors](#)
 - 5.1 [Error &warning messages](#)

"MR-MPI WWW Site"_mws – "MR-MPI Documentation"_md – "OINK Documentation"_od – "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line 1. Building OINK :h3 This section describes how to build and run OINK, which is a simple C++ program which wraps the MapReduce-MPI (MR-MPI) library. 1.1 "Making OINK"_#1_1 1.2 "Building OINK as a library"_#1_2 1.3 "Running OINK"_#1_3 1.4 "Command-line options"_#1_4 all(b) :line 1.1 Making OINK :h4,link(1_1) All of the OINK source files are in the "oink" directory of the MR-MPI distribution tarball. The "src" directory contains the source files for the MR-MPI library itself. These are the 4 steps to building OINK: (1) Insure MPI is installed on your system. (2) Build the MR-MPI library. (3) Use or create a oink/MAKE/Makefile.machine file appropriate for your machine. (4) Type "make machine" :ul Here are more details on each step: (1) MPI installation MPI is the message passing interface library, which is likely already installed on your Linux box or Mac, and on most parallel machines. If not, it is freely available. The two most commonly used generic versions are "OpenMPI"_openmpi and "MPICH"_mpich. Download and install one of these if you need to. The default installation location on Linux is under /usr/local. Or if you do not plan to run the MR-MPI library or OINK in parallel, you can use the provided dummy MPI library in the mpistubs dir. From mpistubs, type "make" and you should get a libmpi.a file. If not, you may need to edit the mpistubs/Makefile. (2) Build the MR-MPI library See "this section"_../doc/Start.html of the MapReduce-MPI library doc pages for instructions on how to do this. When you have done this a file named src/libmrmpi_machine.a should exist. (3) Create a Makefile.machine appropriate for your machine. See the oink/MAKE dir for examples of these. You may be able to use one of these, or edit one that is close to create one for your machine. The only settings you need to worry about are those in the top section. Set the C++ compiler name and settings appropriate for your box. The only extra libraries used by OINK are MPI and MR-MPI. The settings for the latter are already present in the Makefile.machine files. You may need to change the MPI settings depending on how you did your installation. If you use the MPI compiler wrappers (mpiCC) for building an MPI-based program like OINK, then you likely need no additional -I or -L or LIB settings. If you use your system compilers directly, e.g. g++, then you will typically need these MPI-related settings: An -I setting for where to find the file mpi.h. A -L setting for where the MPI library is, libmpich.a A LIB setting for the MPI library, e.g. -lmpich :ul If you are using the provided dummy MPI library (no parallelism), then see MAKE/Makefile.serial for how to compile/link with it. Note that you should insure you build both OINK and the MR-MPI library with the same MPI. If not, confusion will ensue. (4) Type "make machine" Do this from the oink directory where its source files are. If you just type "make" you will see what machine options are available (first line of oink/MAKE/Makefile.machine files). Some other options are also listed, e.g. for cleaning up. If you type "make machine", an executable file oink_machine should be created, e.g. oink_linux or oink_mac. If that happens, you're done. If some error is generated, then you'll need to edit your oink/MAKE/Makefile.machine. Find a local make or machine expert to help if you have problems. If you build OINK on a new kind of machine, for which there isn't a similar Makefile for in the oink/MAKE directory, send it to the developers and we'll add it to the OINK distribution. You can make OINK for multiple platforms from the same oink directory. Each target creates its own object sub-directory called Obj_name where it stores the system-specific *.o files. :line 1.2 Building OINK as a library :h4,link(2_4) OINK can be built as a library, which can then be called from another application or a scripting language. This is done by typing: make makelib make -f Makefile.lib foo :pre where foo is the machine name. The first "make" command will create a current Makefile.lib with all the file names in your src dir. The 2nd "make" command will use it to build OINK as a library. This requires that Makefile.foo have a library target (lib) and system-specific settings for ARCHIVE and ARFLAGS. See Makefile.linux for an example. The build will create the file liboink_foo.a which another application can link to. When used from a C++ program, the library allows one or more OINK objects to be instantiated. All of OINK is wrapped in a OINK_NS namespace; you can safely use any of its classes and methods from within your application code, as needed. When used from a C or Fortran program or a scripting language, the library has a simple C-style interface, provided in oink/library.cpp and oink/library.h. :line 1.3 Running OINK :h4,link(1_3) By default, OINK runs by reading commands from stdin; e.g. oink_linux < in.file. This means you first create an input script (e.g. in.file) containing the desired commands. "This section"_Section_commands.html describes how input scripts are structured and what commands they contain. You can test OINK on any of the sample inputs provided in the examples directory. OINK input scripts are named in.*. Here is how you might run one of the tests on a Linux box, using mpirun to launch a parallel job: cd src

make -f Makefile.linux # builds src/libmrmpi.a cd ../oink make linux # builds oink/oink_linux cd ../examples

`mpirun -np 4 ../oink/oink_linux ../doc/*.txt < in.wordcount` :pre If OINK encounters errors in the input script or while running a command it will print an ERROR message and stop or a WARNING message and continue. See "this section" [_Section_errors.html](#) for a discussion of the various kinds of errors OINK can or can't detect, a list of all ERROR and WARNING messages, and what to do about them. OINK can run a MapReduce calculation on any number of processors, including a single processor. :line 2.6 Command-line options :h4,link(2_6)

At run time, OINK recognizes several optional command-line switches which may be used in any order. Either the full word or the one-letter abbreviation can be used: `-echo` or `-e` `-partition` or `-p` `-in` or `-i` `-log` or `-l` `-screen` or `-s` `-var` or `-v` :ul For example, `oink_ibm` might be launched as follows: `mpirun -np 16 oink_ibm -var file tmp.out -log my.log -screen none < in.graph` :pre Here are the details on the options: `-echo style` :pre Set the style of command echoing. The style can be {none} or {screen} or {log} or {both}. Depending on the style, each command read from the input script will be echoed to the screen and/or logfile. This can be useful to figure out which line of your script is causing an input error. The default value is {log}. The echo style can also be set by using the "echo" [_echo.html](#) command in the input script itself. `-partition 8x2 4 5 ...` :pre Invoke OINK in multi-partition mode. When OINK is run on P processors and this switch is not used, OINK runs in one partition, i.e. all P processors run a single calculation. If this switch is used, the P processors are split into separate partitions and each partition runs its own calculation. The arguments to the switch specify the number of processors in each partition. Arguments of the form MxN mean M partitions, each with N processors. Arguments of the form N mean a single partition with N processors. The sum of processors in all partitions must equal P. Thus the command "`-partition 8x2 4 5`" has 10 partitions and runs on a total of 25 processors. Note that with MPI installed on a machine (e.g. your desktop), you can run on more (virtual) processors than you have physical processors. The input script specifies what simulation is run on which partition; see the "variable" [_variable.html](#) and "next" [_next.html](#) commands. `-in file` :pre Specify a file to use as an input script. This is an optional switch when running OINK in one-partition mode. If it is not specified, OINK reads its input script from stdin – e.g. `oink_linux < in.run`. This is a required switch when running OINK in multi-partition mode, since multiple processors cannot all read from stdin. `-log file` :pre Specify a log file for OINK to write status information to. In one-partition mode, if the switch is not used, OINK writes to the file `log.oink`. If this switch is used, OINK writes to the specified file. In multi-partition mode, if the switch is not used, a `log.oink` file is created with hi-level status information. Each partition also writes to a `log.oink.N` file where N is the partition ID. If the switch is specified in multi-partition mode, the hi-level logfile is named "file" and each partition also logs information to a `file.N`. For both one-partition and multi-partition mode, if the specified file is "none", then no log files are created. Using a "log" [_log.html](#) command in the input script will override this setting. `-screen file` :pre Specify a file for OINK to write its screen information to. In one-partition mode, if the switch is not used, OINK writes to the screen. If this switch is used, OINK writes to the specified file instead and you will see no screen output. In multi-partition mode, if the switch is not used, hi-level status information is written to the screen. Each partition also writes to a `screen.N` file where N is the partition ID. If the switch is specified in multi-partition mode, the hi-level screen dump is named "file" and each partition also writes screen information to a `file.N`. For both one-partition and multi-partition mode, if the specified file is "none", then no screen output is performed. `-var name value1 value2 ...` :pre Specify a variable that will be defined for substitution purposes when the input script is read. "Name" is the variable name which can be a single character (referenced as \$x in the input script) or a full string (referenced as \${abc}). An "index-style variable" [_variable.html](#) will be created and populated with the subsequent values, e.g. a set of filenames. Using this command-line option is equivalent to putting the line "variable name index value1 value2 ..." at the beginning of the input script. Defining an index variable as a command-line argument overrides any setting for the same index variable in the input script, since index variables cannot be re-defined. See the "variable" [_variable.html](#) command for more info on defining index and other kinds of variables and "this section" [_Section_commands.html#3_2](#) for more info on using variables in input scripts. "MR-MPI WWW Site" [_mws](#) – "MR-MPI Documentation" [_md](#) – "OINK Documentation" [_od](#) – "OINK Commands" [_oc](#) :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,../doc/Manual.html) :link(od,Manual.html) :link(oc,Section_script.html#comm)

:line 2. OINK Commands :h3 This section describes OINK input scripts and what commands are used to define an OINK calculation. 2.1 "Input script operation" [_#2_1](#) 2.2 "Parsing rules" [_#2_2](#) 2.3 "Input script commands" [_#2_3](#) :all(b) :line 2.1 Input script operation :link(3_1),h4 OINK executes by reading commands from a input script (text file), one line at a time.

When the input script ends, OINK exits. Each command causes OINK to take some action. It may set an internal variable, read in a file, or perform a MapReduce operation. Most commands have default settings, which means you only need to use the command if you wish to change the default. Note that OINK does not read your entire input script and then perform a calculation with all the settings. Rather, the input script is read one line at a time and each command takes effect when it is read. Thus this sequence of commands: `set verbosity 1 mr foo :pre` does something different than this sequence: `mr foo set verbosity 1 :pre` In the first case, the MR object created will have its verbosity set to 1. In the latter case it will have the default verbosity of 0, since the set command was not used until after the MR object was created. Many input script errors are detected by OINK and an ERROR or WARNING message is printed. "This section" [_Section_errors.html](#) gives more information on what errors mean. The documentation for each command gives additional information. [:line 2.2 Parsing rules :link\(3_2\),h4](#) Each non-blank line in the input script is treated as a command. OINK commands are case sensitive. Pre-defined command names are lower-case, as are specified command arguments. Upper case letters may be used in file names or user-chosen ID strings. Here is how each line in the input script is parsed by OINK: (1) If the last printable character on the line is a `"` character (with no surrounding quotes), the command is assumed to continue on the next line. The next line is concatenated to the previous line by removing the `"` character and newline. This allows long commands to be continued across two or more lines. (2) All characters from the first `"#"` character onward are treated as comment and discarded. See an exception in (6). Note that a comment after a trailing `"` character will prevent the command from continuing on the next line. Also note that for multi-line commands a single leading `"#"` will comment out the entire command. (3) The line is searched repeatedly for `$` characters, which indicate variables that are replaced with a text string. See an exception in (6). If the `$` is followed by curly brackets, then the variable name is the text inside the curly brackets. If no curly brackets follow the `$`, then the variable name is the single character immediately following the `$`. Thus `${myTemp}` and `$x` refer to variable names "myTemp" and "x". See the "variable" [_variable.html](#) command for details of how strings are assigned to variables and how they are substituted for in input script commands. (4) The line is broken into "words" separated by whitespace (tabs, spaces). Note that words can thus contain letters, digits, underscores, or punctuation characters. (5) The first word is the command name. All successive words in the line are arguments. (6) If you want text with spaces to be treated as a single argument, it can be enclosed in either double or single quotes. E.g. `print "Value = $t" print 'Value = $t' :pre` The quotes are removed when the single argument is stored internally. See the "if" [_if.html](#) commands for examples. A `"#"` or `"$"` character that is between quotes will not be treated as a comment indicator in (2) or substituted for as a variable in (3). **IMPORTANT NOTE:** If the argument is itself a command that requires a quoted argument (e.g. using a "print" [_print.html](#) command as part of an "if" [_if.html](#) command), then the double and single quotes can be nested in the usual manner. See the doc pages for those commands for examples. Only one of level of nesting is allowed, but that should be sufficient for most use cases.

[:line 2.3 Input script commands :h4,link\(2_3\)](#) There are 4 kinds of OINK commands: (1) Set command to alter parameters: "set" [_set.html](#) [:ul](#) (2) MR-MPI library commands: "mr foo" [_mr.html](#) "foo map ..., foo reduce ..., etc" [_mrmapi.html](#) [:ul](#) (3) Named commands: "input" [_input.html](#) "output" [_output.html](#) "myfoo params ... -i ... -o ..." [_command.html](#) [:ul](#) (4) Miscellaneous commands that are part of the scripting language: "clear" [_clear.html](#) "echo" [_echo.html](#) "if" [_if.html](#) "include" [_include.html](#) "jump" [_jump.html](#) "label" [_label.html](#) "log" [_log.html](#) "next" [_next.html](#) "print" [_print.html](#) "shell" [_shell.html](#) "variable" [_variable.html](#) [:ul](#) [:line :link\(comm\)](#) Here is a list of all OINK input script commands alphabetically: "clear" [_clear.html](#), "echo" [_echo.html](#), "if" [_if.html](#), "include" [_include.html](#), "input" [_input.html](#), "jump" [_jump.html](#), "label" [_label.html](#), "log" [_log.html](#), "mr" [_mr.html](#), "MR-MPI library commands" [_mrmapi.html](#), "named commands" [_command.html](#), "next" [_next.html](#), "output" [_output.html](#), "print" [_print.html](#), "set" [_set.html](#), "shell" [_shell.html](#), "variable" [_variable.html](#) [:tb\(c=6,ea=c\)](#) These are the named commands currently included in OINK. We will add to this list from time to time. If you write a useful new command, send it to us and we can include it in the distribution. "degree" [_degree.html](#), "neigh_tri" [_neigh_tri.html](#), "neighbor" [_neighbor.html](#), "rmat" [_rmat.html](#), "sgi_enumerate" [_sgi_enumerate.html](#), "sgi_prune" [_sgi_prune.html](#), "sgi_sample" [_sgi_sample.html](#), "tri_find" [_tri_find.html](#), "wordfreq" [_wordfreq.html](#) [:tb\(c=2,ea=c\)](#) "MR-MPI WWW Site" [_mws](#) - "MR-MPI Documentation" [_md](#) - "OINK Documentation" [_od](#) - "OINK Commands" [_oc](#) [:c](#) [:link\(mws,http://www.sandia.gov/~sjplimp/mapreduce.html\)](#) [:link\(md,./doc/Manual.html\)](#) [:link\(od,Manual.html\)](#) [:link\(oc,Section_script.html#comm\)](#) [:line 3.](#) Adding Callback Functions to OINK [:h3](#) In the oink directory, the files `map_*.cpp`, `reduce_*.cpp`, `compare_*.cpp`, `hash_*.cpp`, and `scan_*.cpp` each contain one or more functions

which can be used as callback methods, passed to MR–MPI library calls, such as the "map()"_doc/.map.html and "reduce()"_./doc/reduce.html operations. This can be done either in "named commands"_command.html that you write, as described in "this section"_Section_commands.html of the documentation, or in "MR–MPI library commands"_mrmmpi.html made directly from an OINK input script. The collection of these files and callback functions is effectively a library of tools that can be used by new "named commands"_command.html or your input script to speed the development of new MapReduce algorithms and workflows. Over time, we intend to add new callback function to OINK, and also invite users to send their own functions to the developers for inclusion in OINK. The map(), reduce(), and scan() callback functions include a "void *ptr" as a final argument, which the caller can pass to the callback. This is typically done to enable the callback function to access additional parameters stored by the caller. When doing this with functions listed in the map_*.cpp, reduce_*.cpp, and scan_*.cpp files in OINK, you will want to make the data these pointers point to "portable", so that and "named command" can use it. Thus you would should not typically encode class–specific or command–specific data in the structure pointed to. Instead, your caller should create the minimal data structure that the callback function needs to operate, and store the structure in a map_*.h file that corresponds to the specific map_*.cpp file that contains the function (or reduce_*.h or scan_*.h). See the file oink/map_rmat.h file as an example. It contains the definition of an RMat_params structure, which is used by both the "rmat command"_rmat.txt and the map() methods it uses, listed in map_rmat.cpp. Both the rmap.cpp and map_rmat.cpp files include the map_rmat.h header file to accomplish this. Other commands or callback functions could use the same data structure by including that header file. The following sections list the various callback function currently included in OINK, and a brief explanation of what each of them does. Note that map() functions come in 4 flavors, depending on what MR–MPI library "map() method"_./doc/map.html is being used. Similarly, scan() functions come in 2 flavors, as documented on the "scan() method"_./doc/scan.html page. Map_*.cpp and scan_*.cpp files within OINK can contain any of the 4 or 2 flavors of map() and scan() methods.

3.1 "Map() functions"_{3_1} 3.2 "Reduce() functions"_{3_2} 3.3 "Compare() functions"_{3_3} 3.4 "Hash() functions"_{3_4} 3.5 "Scan() functions"_{3_5}

:all(b) The documenation below this double line is auto–generated when the OINK manual is created. This is done by extracting C–style documentation text from the map_*.cpp, reduce_*.cpp, compare_*.cpp, hash_*.cpp, and scan_*.cpp files in the oink directory. Thus you should not edit content below this double line. In the *.cpp files in the oink directory, the lines between a line with a "/*" and a line with a "*/" are extracted. In the tables below, the first such line of extracted text is assumed to be the function name and appears in the left column. The remaining lines appear in the right columns.

:line :line Map() functions

:link(3_1),h4

rmat_generate	generate RMat matrix entries emit one KV per edge: key = edge, value = NULL
rmat_generate	generate RMat matrix entries emit one KV per edge: key = edge, value = NULL
rmat_stats	

	print # of rows with a specific # of nonzeroes
--	------------------------------------------------------------

:line Reduce() functions :link(3_1),h4

cull	eliminate duplicate edges input: one KMV per edge, MV has multiple entries if duplicates exist output: one KV per edge: key = edge, value = NULL
degree	count nonzeroes in each row input: one KMV per row, MV has entry for each nonzero output: one KV: key = # of nonzeroes, value = NULL
histo	count rows with same # of nonzeroes input: one KMV per nonzero count, MV has entry for each row output: one

	KV: key = # of nonzeroes, value = # of rows
nonzero	enumerate nonzeroes in each row input: one KMV per edge output: one KV per edge: key = row I, value = NULL
sum_count	compute count from nvalues input: one KMV per edge, MV has multiple entries if duplicates exist output: one KV per edge: key = edge, value = NULL

:line Compare() functions :link(3_1),h4

:line Hash() functions :link(3_1),h4

:line Scan() functions :link(3_1),h4

print_edge	print out an edge input: one KMV per edge, MV has multiple entries if duplicates exist output: one KV per edge:
------------	-----------------------------------------------------------------------------------------------------------------------------------------------------

	key = edge, value = NULL
print_string_int	print out string and int input: one KMV per edge, MV has multiple entries if duplicates exist output: one KV per edge: key = edge, value = NULL

:line "MR-MPI WWW Site"_mws – "MR-MPI Documentation"_md – "OINK Documentation"_od – "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..../doc/Manual.html) :link(od,Manual.html) :link(oc,Section_script.html#comm) :line 4. Adding Commands to OINK :h3 The purpose of this section is to give details of how to write new "named commands"_command.html that can be added to OINK and which will be invocable by your input scripts and will interact appropriately with other OINK commands. OINK is designed to make this easy to do with a minimum of special coding on your part. Several such named commands are included with the OINK distribution; more will be added over time. See "this section"_Section_script.html#comm of the manual for a list of the current named commands in OINK. We also invite OINK users to send email to the developers with new commands they have written and wish to share, so we can add them to the distribution, attributed to you. 4.1 "Source files for the new class"_#4_1 4.2 "Methods in the new class"_#4_2 4.3 "Calls to the OINK object manager"_#4_3 4.4 "Calling back to map() and reduce() functions"_#4_4 :all(b) :line :line 4.1 Source files for the new class :link(4_1),h4 In OINK a named command is a child class that derives from the Command parent class (see src/command.cpp and src/command.h), meaning that it contains several methods that can be called by the OINK framework. Adding a new named command to OINK is as simple as writing the code for it in two new files (e.g. foo.cpp and foo.h), dropping them into the src directory, and re-building OINK. It is easiest to understand the description that follows if you look at an example named command in OINK. In what follows we will use the "degree"_dere.html command, contained in src/degree.cpp and src/degree.h for illustration purposes. The *.h file for a new named command should have lines like these at the top (from the src/degree.h file): #ifdef COMMAND_CLASS CommandStyle(degree,Degree) #else :pre CommandStyle(arg1,arg2) is a macro that gets converted by the OINK build procedure into source code. Arg1 is the "name" of the named command, which is how you reference it in your input script, e.g. as degree -i graphdir -o out/outfile NULL :pre Arg2 is the class that implements that command. The list of all such named commands will appear in the style_command.h file after OINK is (re)built, via a make command. The remainder of the *.h file (between the #else and final pair of #endif) is the definition of your new class. Note that you will need to include the mapreduce.h file (from the MR-MPI library src dir) and the MAPREDUCE_NS namespace if your class definition includes any map(), reduce(), etc callback functions since they have the "KeyValue" class name in their prototype. :line :line 4.2 Methods in the new class :link(4_2),h4 When a line in the input script starts with a "named command"_command.html, the associated class is instantiated, its params() and run() methods are called, and the instance of the class is destroyed. The constructor of your new class should set two variables, ninputs and noutputs, which are the number of input and

output descriptors it requires you to list in the input script. For the "degree"_degree.html command, ninputs = noutpus = 1, as illustrated in the example above. Your new class is required to define only two methods: params() and run(). Params() is passed the list of arguments following the command name, excluding the "-i" and "-o" arguments, which are processed separately by the parent Command class. Your params() method should parse and check the arguments and generate an error message if the number of arguments is incorrect or any of their values is invalid. Note that the "degree"_degree.html command takes 0 arguments. Run() is called to invoke the meat of your command and it can perform any series of MapReduce or other operations you wish, using as many "MapReduce objects"_md (from the MR-MPI library) and "MR-MPI objects"_mr.html (managed by OINK) as you wish. The calls that run() can make to the OINK object manager (obj) are discussed in the next section. The destructor of your new class should free any memory it has allocated, including any local MapReduce objects that it allocated. Note that this is different from "MR-MPI objects"_mr.html and the underlying MapReduce objects they wrap, and which are often associated with the input and output descriptors of your command in the input script; those objects are created/destroyed by OINK itself, as discussed in the next section.

:line :line 4.3

Calls to the OINK object manager :link(4_3),h4 These are the calls that the run() method of your new class can make to the OINK object manager. Each is discussed below. Note that you make the calls via the "obj" pointer which is visible to your class, e.g. obj->cleanup(). This means you should add the line #include "object.h" :pre at the top of your *.cpp file.

```
MapReduce *Object::create_mr();
MapReduce *Object::create_mr(int verbosity, int timer, int memsize, int outofcore);
MapReduce *Object::copy_mr(MapReduce *mr);
int Object::permanent(MapReduce *mr);
:pre MapReduce *Object::input(int index);
MapReduce *Object::input(int index, void (*map1)(int, char *, KeyValue *, void *), void *ptr);
MapReduce *Object::input(int index, void (*map2)(int, char *, int, KeyValue *, void *), void *ptr);
MapReduce *Object::input(int index, void (*map1)(int, char *, KeyValue *, void *), void (*map2)(int, char *, int, KeyValue *, void *), void *ptr);
:pre void Object::output(int index, MapReduce *mr);
void Object::output(int index, MapReduce *mr, void (*scankv)(char *, int, char *, int, void *), void *ptr, int disallow);
void Object::output(int index, MapReduce *mr, void (*scankmv)(char *, int, char *, int, int *, void *), void *ptr, int disallow);
void Object::output(int index, MapReduce *mr, void (*map)(uint64_t, char *, int, char *, int, KeyValue *, void *), void *ptr, int disallow);
void Object::output(int index, MapReduce *mr, void (*reduce)(char *, int, char *, int, int *, KeyValue *, void *), void *ptr, int disallow);
void Object::output(int index, MapReduce *mr, void (*scankv)(char *, int, char *, int, void *), void (*scankmv)(char *, int, char *, int, int *, void *), void (*map)(uint64_t, char *, int, char *, int, KeyValue *, void *), void (*reduce)(char *, int, char *, int, int *, KeyValue *, void *), void *ptr, int disallow);
:pre void Object::cleanup();
:pre
```

Here is a brief summary of the calls your run() method will typically make: 1 call to input() for each of its Ninput input descriptors calls to create_mr() or copy_mr() for additional MapReduce objects it uses 1 call to output() for each of its Noutput output descriptors final call to cleanup() at the end

:ul The details are discussed below.

:line Your run() method should call one of the 4 variants of the input() methods one time for each of its inputs. Which variant it calls depends on what forms of input you wish to support, which is related to the "-i" arguments specified with your "named command"_command.html in the input script and additional options set by the "input"_input.html command. Each call takes an "index" argument which is the index of the input descriptor being referenced, from 1 to Ninputs. Each call returns a pointer to a MapReduce object which will contain the desired input data as key/value (KV) pairs. As the "named command"_command.html doc page explains, each input descriptor for your command can be specified in the input script as one or more files or directories or as an existing MR-MPI object. For reading files, there are 2 kinds of map() methods that can be used to convert the file contents into KV pairs, one where a filename is passed to your callback function, and the other where a chunk of bytes is passed to your callback function. See the "map() method"_.doc/map.html doc page for details. If you invoke this method: MapReduce *Object::input(int index); :pre then the input descriptor must be specified in your input script as an existing MR-MPI object. No reading of files is allowed. If you invoke this method: MapReduce *Object::input(int index, void (*map1)(int, char *, KeyValue *, void *), void *ptr); :pre then the input can be specified as either an MR-MPI object or as files which will be processed via the map1() callback function which receives a filename as an argument, so that it can open the file, read it, and generate KV pairs. If you invoke this method: MapReduce *Object::input(int index, void (*map2)(int, char *, int, KeyValue *, void *), void *ptr); :pre then the input can be specied as either an MR-MPI object or as files which will be processed via the map2() callback function which receives a chunk of bytes read from a file as an arbument, so that it can convert the byte string into KV pairs. To use this map2() method, you would also need to specify an

"input"_input.html command in your input script that setup various options needed to call the "MR-MPI library map() method" `../doc/map.html` that uses `map2()` as a callback function. If you invoke the 4th variant: `MapReduce *Object::input(int index, void (*map1)(int, char *, KeyValue *, void *), void (*map2)(int, char *, int, KeyValue *, void *), void *ptr);` :pre then both kinds of `map()` callback functions can be specified, `map1()` and `map2()`, and OINK will select which to use depending on what options have been setup via the "input"_input.html command for this input descriptor. Note that you have to provide the `map1()` and/or `map2()` callback functions to the `input()` calls, with the correct prototype. As discussed below and on "this doc page" `_Section_functions.html`, they can be static methods in your class, or they can be `map()` methods in separate files in the OINK src directory, which are named `map_*.cpp`. Also note that if you want to provide maximum flexibility for using your "named command"_command.html, then you should provide one of both flavors of callback `map()` functions for allowing input from files along with input from an existing MR-MPI object. If you do not provide either callback or just one of the two, then input scripts will be limited in what forms of input descriptor they can define. :line Your `run()` method should call one of the 6 variants of the `output()` methods one time for each of its outputs. Which variant it calls depends on what forms of output you wish to support, which is related to the "-o" arguments specified with your "named command"_command.html in the input script and additional options set by the "output"_output.html command. Each call takes an "index" argument which is the index of the output descriptor being referenced, from 1 to `Noutputs`. Each call also takes a `MapReduce` object pointer "mr", which contains the data you wish to output. As the "named command"_command.html doc page explains, each output descriptor for your command is specified in the input script with 2 parts, either of which can be NULL. The first part is a filename for writing output to files. The second part is the ID of an MR-MPI object which will contain the output. For writing files, there are 4 kinds of callback methods that can be used to write the contents of "mr" to a file. Each of these 4 methods is called with a "FILE *" as its final "void *" argument. This is the file pointer to a file created and opened (and later closed) by OINK which the callback method can write its data to. If you pass your own non-NULL pointer to the callback method via the "void *ptr" argument to the `output()` calls, then it will be appended to the FILE *, so that it can be dereferenced as a 2nd pointer passed to the callback function. If you invoke this method: `void Object::output(int index, MapReduce *mr);` :pre then the output descriptor must be specified in your input script as only defining a MR-MPI object for output. No writing to files is allowed. This call will assign the ID specified in your input script to the MR-MPI object that wraps "mr". Also note, that this will remove the ID from any other MR-MPI object that has the same ID. They then become unnamed or temporary MR-MPI objects which will be deleted at the end of your `run()` method. See further discussion of temporary versus permanent MR-MPI objects in the next section. If you invoke one of these 2 methods: `void Object::output(int index, MapReduce *mr, void (*scankv)(char *, int, char *, int, void *), void *ptr, int disallow);` `void Object::output(int index, MapReduce *mr, void (*scankmv)(char *, int, char *, int, int *, void *), void *ptr, int disallow);` :pre then the output can be specified as either an MR-MPI object or as files which will be written to via the `scankv()` or `scankmv()` callback functions respectively. In the first case, the `scankv()` function will receive key/value (KV) pairs, one at a time from the "mr" `MapReduce` object. In the second case, the `scankmv()` function will receive key/multivalue (KMV) pairs, one at a time from the "mr" `MapReduce` object. The `MapReduce` object will be unaltered by this operation. See the "scan() method" `../doc/scan.html` doc page in the MR-MPI library for details. The "disallow" flag is explained below. If you invoke one of these 2 methods: `void Object::output(int index, MapReduce *mr, void (*map)(uint64_t, char *, int, char *, int, KeyValue *, void *), void *ptr, int disallow);` `void Object::output(int index, MapReduce *mr, void (*reduce)(char *, int, char *, int, int *, KeyValue *, void *), void *ptr, int disallow);` :pre then the output can be specified as either an MR-MPI object or as files which will be written to via the `map()` or `reduce()` callback functions respectively. In the first case, the `map()` function will receive key/value (KV) pairs, one at a time from the "mr" `MapReduce` object. In the second case, the `reduce()` function will receive key/multivalue (KMV) pairs, one at a time from the "mr" `MapReduce` object. For the first case, the `MapReduce` object will typically be unaltered by this operation, since the "MR-MPI library `map()` method" `../doc/map.html` is called with `addflag=1`, so that the existing KV pairs are preserved. But your `map()` callback function should not emit any new KV pairs. For the second case, the `MapReduce` object will be altered by this operation, since the "MR-MPI library `reduce()` method" `../doc/reduce.html` deletes the KMV pairs and replaces them with new KV pairs which your `reduce()` callback function may or may not emit. The "disallow" flag is explained below. If you invoke the 6th variant: `void Object::output(int index, MapReduce *mr, void (*scankv)(char *, int, char *, int, void *), void (*scankmv)(char *, int, char *, int, int *, void *), void`

(`*map`)(`uint64_t`, `char *`, `int`, `char *`, `int`, `KeyValue *`, `void *`), `void (*reduce)`(`char *`, `int`, `char *`, `int`, `int *`, `KeyValue *`, `void *`), `void *ptr`, `int disallow`); :pre then any of the 4 kinds of callback functions can be specified, namely `scankv()`, `scankmv()`, `map()`, or `reduce()`. Those that you do not wish to provide or that are not compatible with the current state of the MapReduce object "mr" (which will contain either kV or KMV pairs, but not both), can be specified as NULL. Note that you have to provide each of these 4 callback functions to the `output()` calls, with the correct prototype. As discussed below and on "this doc page"_Section_functions.html, they can be static methods in your class, or they can be methods in separate files in the OINK src directory, which are named `scan_*.cpp`, `map_*.cpp`, and `reduce_*.cpp` respectively. Also note that if you want to provide maximum flexibility for using your "named command"_command.html, then you should provide at least one flavor of a callback function for allowing output to files along with output to an MR-MPI object. If you do not do this, then input scripts will be limited in what forms of output descriptor they can define. All but the first of the `output()` variants can be called with an optional `disallow` flag which is set to 0 by default. If these methods are called with `disallow=1`, then no output to an MR-MPI object is allowed. This is useful if you expect the `run()` method of your "named command"_command.html to subsequently change the data stored in the MapReduce object, and thus make the data written to an output file differ from what is stored in the MapReduce object. :line Your `run()` method may need to use additional MapReduce objects as workspace, in addition to its inputs. Some of these may end up holding the data you wish to output. One key point to understand is that the OINK object manager keeps track of two kinds of MR-MPI objects, each of which is a thin wrapper on MapReduce objects which hold your key/value (KV) or key/multivalue (KMV) data. Each MR-MPI object can be "permanent" meaning it has an ID which can be referenced by input script commands. Or it can be "temporary", meaning it has no ID and was created to hold data input from a file or by the function calls discussed below. Permanent MR-MPI objects persist until they are explicitly deleted by your input script. Temporary MR-MPI objects are deleted at the end of your `run()` method; they can be thought of as workspace created and used by your `run()` method. These two calls create a new temporary MR-MPI object and return a pointer to the MapReduce object contained within it: `MapReduce *Object::create_mr()`; `MapReduce *Object::create_mr(int verbosity, int timer, int memsize, int outofcore)`; :pre The first variant will use the default settings for the MapReduce object; see the "set"_set.html command and the "settings doc page"../doc/settings.html of the MR-MPI library for details. The second variant allows you to override a few of the settings with specified values. This call makes a copy of the "mr" MapReduce object, wraps it in a new temporary MR-MPI object, and returns a pointer to the new MapReduce object: `MapReduce *Object::copy_mr(MapReduce *mr)`; :pre There are two reasons to create new and copied MapReduce objects via these calls, rather than directly invoking MR-MPI library calls within your `run()` method, e.g. do this `MapReduce *mr = obj->create_mr()`; `MapReduce *mr2 = obj->copy(mr)`; instead of this: `MapReduce *mr = new MapReduce()`; `MapReduce *mr2 = mr->copy()`; The first reason is that OINK will now manage the memory associated with the new MR-MPI objects and free them for you at the end of your `run()` method; see the `cleanup()` method discussion below. The second is that you can assign an ID to these temporary MR-MPI objects via the `output()` calls discussed above, which you cannot do if you create the MapReduce object yourself. I.e. you cannot pass a pointer to `output()` of a MapReduce object that you created yourself if that operation will assign an ID to the MR-MPI object that was specified in your input script. You are of course free to create additional MapReduce objects yourself via direct calls to the MR-MPI library. In this case you should insure you free the objects explicitly before the `run()` method ends, so as not to leak memory. This call is useful for checking whether a MR-MPI object has been assigned a name or not, when it was used for input or output: `int Object::permanent(MapReduce *mr)`; :pre It is called using a MapReduce object pointer and returns a 1 if the associated MR-MPI object that wraps it has a name, and a 0 if it does not. There are two uses for this call. First, it can be used after an `input()` call to determine whether the input was done from a file or an existing MR-MPI object. In the former case `permanent()` will return 0, since the new MR-MPI object holding the data is unnamed. In the latter case it will return 1, since the MR-MPI object holding the data was named in your input script as one of the "-i" arguments to the "named command"_command.html. If the `run()` method will subsequently alter the MapReduce object and it is permanent, you can make a copy of it, so as to not alter the original. Second, it can be used after an `output()` call to determine whether the MapReduce object was assigned a name. This will be the case if a MR-MPI ID was specified in your input script as one of the "-o" arguments to the "named command"_command.html. If this is the case, you typically do not want to alter the data in the MapReduce object after outputting it. If you wish to further process the data, you can make a copy. Finally, this method should be

called at the end of your run() method to free all the temporary MR-MPI objects stored by OINK, and perform other internal cleanup: void Object::cleanup(); :pre :line 4.4 Calling back to map() and reduce() functions

:link(4_4),h4 You run() method will typically invoke various methods from the MR-MPI library which involve callback functions, e.g. for performing map() or reduce() operations. The "MR-MPI library manual" [_./doc/Technical.html#callback](http://www.sandia.gov/~sjplimp/mapreduce.html#callback) discusses the general rules for passing a pointer to a callback function to a MR-MPI library method. Since you will be doing this from within the class that encodes your "named command" [_command.html](http://www.sandia.gov/~sjplimp/mapreduce.html#command) you have two choices. First, you can pass a pointer to a static function declared within your class. This function cannot directly access any class variables, but you can pass it the "this" pointer for the class (as the void * argument to the map() or reduce() function) which the callback function can use to access class variables indirectly, through that pointer. If you do this, then the map() and reduce() methods defined in your class can only be used by that "named command" [_command.html](http://www.sandia.gov/~sjplimp/mapreduce.html#command). An alternative is to put your callback functions in their own files, named map_*.cpp for map() functions, reduce_*.cpp for reduce() functions, compare_*.cpp for compare() functions, hash_*.cpp for hash() functions, and scan_*.cpp functions. By doing this the callback functions can be used by any "named command" [_command.html](http://www.sandia.gov/~sjplimp/mapreduce.html#command) or as arguments to the "MR-MPI library commands" used in an input script to invoke the MR-MPI library methods directly. See the oink/rmat.cpp file, which implements the "rmat" [_rmat.html](http://www.sandia.gov/~sjplimp/mapreduce.html#rmat) command, for an example of a "named command" [_command.html](http://www.sandia.gov/~sjplimp/mapreduce.html#command) which accesses several of its callback functions in this manner. Each map_*.cpp file (and reduce_*.cpp, compare_*.cpp, etc) can contain one or more map() (reduce(), compare(), etc) callback functions. These are not class methods, but stand-alone functions. See examples in the oink directory. The header files that contain the prototypes for these functions are named style_map.h, style_reduce.h, etc and are auto-generated when OINK is built. Your "named command" [_command.html](http://www.sandia.gov/~sjplimp/mapreduce.html#command) class, e.g. rmat.cpp, simply needs to include these style header files in order to use any of the callback functions in OINK. Likewise, any callback function included in one of these files can be accessed by name in your input script when using one of the "MR-MPI library commands" [_mrmpi.html](http://www.sandia.gov/~sjplimp/mapreduce.html#mrmpi). Documentation for the collection of map(), reduce(), etc functions is also auto-extracted and included in "this section" [_Section_functions.html](http://www.sandia.gov/~sjplimp/mapreduce.html#section_functions) of the OINK documentation. Instructions on how to pass generic pointers to the callback functions is also discussed in "this section" [_Section_functions.html](http://www.sandia.gov/~sjplimp/mapreduce.html#section_functions). It is also possible in the run() method of your "named command" [_command.html](http://www.sandia.gov/~sjplimp/mapreduce.html#command) to select a callback function based on an input script parameter to your command. For example, the input script could list the name of a particular compare() function you wish to use to sort the data in a MapReduce object. By calling the appropriate lookup() method in the MRMPI class (oink/mrpmi.cpp), the parameter string can be converted into a matching function pointer. For example, consider these lines of code: MapReduce *mr = obj->create_mr(); ... CompareFnPtr compare = compare_lookup(userparam); mr->sort_keys(compare); :pre In this example "userparam" is a string, listed in the input script as a command parameter, which contains a function name, e.g. mySpecialCompare. Assuming that function is included in OINK in a compare_*.cpp file, the the compare_lookup() method will be able to match the string to the function and return a pointer to the function which can then be used as an argument to the "sort_keys()" [_./doc/sort.html](http://www.sandia.gov/~sjplimp/mapreduce.html#sort_keys) MR-MPI library method. The definition of CompareFnPtr and all other callback function pointers is in the "typedefs.h" file, which can be included at the top of your "named command" [_command.html](http://www.sandia.gov/~sjplimp/mapreduce.html#command) *.cpp file. "MR-MPI WWW Site" [_mws](http://www.sandia.gov/~sjplimp/mapreduce.html#mws) - "MR-MPI Documentation" [_md](http://www.sandia.gov/~sjplimp/mapreduce.html#md) - "OINK Documentation" [_od](http://www.sandia.gov/~sjplimp/mapreduce.html#od) - "OINK Commands" [_oc](http://www.sandia.gov/~sjplimp/mapreduce.html#oc) :c

:link(mws,<http://www.sandia.gov/~sjplimp/mapreduce.html>) :link(md,[_./doc/Manual.html](http://www.sandia.gov/~sjplimp/mapreduce.html#md)) :link(od,[_./doc/Manual.html](http://www.sandia.gov/~sjplimp/mapreduce.html#od)) :link(oc,[_./doc/Section_commands.html#comm](http://www.sandia.gov/~sjplimp/mapreduce.html#oc)) :line 5. Errors :h3 OINK tries to flag errors and print informative error messages so you can fix the problem. If you get an error message about an invalid command in your input script, you can determine what command is causing the problem by looking in the log.oink file or using the "echo command" [_echo.html](http://www.sandia.gov/~sjplimp/mapreduce.html#echo) to see it on the screen. For a given command, OINK expects certain arguments in a specified order. If you mess this up, OINK will often flag the error. Generally, OINK will print a message to the screen and logfile and exit gracefully when it encounters a fatal error. Sometimes it will print a WARNING to the screen and logfile and continue on; you can decide if the WARNING is important or not. If OINK crashes or hangs without spitting out an error message first then it could be a bug If you think you have found a bug in OINK, please send an email to the "developers" [_http://www.sandia.gov/~sjplimp/mapreduce.html](http://www.sandia.gov/~sjplimp/mapreduce.html#developers) with info about the problem. Anything you can do to isolate the problem and reproduce it on a small data set will be helpful. :line 5.1 Error & warning messages :h4,link(5_1) These are two alphabetic lists of the "ERROR" [_#error](http://www.sandia.gov/~sjplimp/mapreduce.html#error) and "WARNING" [_#warn](http://www.sandia.gov/~sjplimp/mapreduce.html#warn) messages OINK prints out and the reason why. If the explanation here is not sufficient, the

documentation for the offending command may help. Grepping the source files for the text of the error message and staring at the source code and comments is also not a bad idea! Note that sometimes the same message can be printed from multiple places in the code. Errors: :h4,link(error) :dlb {All universe/uloop variables must have same # of values} :dt Self-explanatory. :dd {All variables in next command must be same style} :dt Self-explanatory. :dd Cannot attempt to open a 2nd input script, when the original file is still being processed. :dd {Arccos of invalid value in variable formula} :dt Argument of arccos() must be between -1 and 1. :dd {Arcsin of invalid value in variable formula} :dt Argument of arcsin() must be between -1 and 1. :dd {Cannot open input script %s} :dt Self-explanatory. :dd {Cannot open log.oink} :dt The default OINK log file cannot be opened. Check that the directory you are running in allows for files to be created. :dd {Cannot open logfile %s} :dt The OINK log file specified in the input script cannot be opened. Check that the path and name are correct. :dd {Cannot open logfile} :dt The OINK log file named in a command-line argument cannot be opened. Check that the path and name are correct. :dd {Cannot open screen file} :dt The screen file specified as a command-line argument cannot be opened. Check that the directory you are running in allows for files to be created. :dd {Cannot open universe log file} :dt For a multi-partition run, the master log file cannot be opened. Check that the directory you are running in allows for files to be created. :dd {Cannot open universe screen file} :dt For a multi-partition run, the master screen file cannot be opened. Check that the directory you are running in allows for files to be created. :dd {Cannot redefine variable as a different style} :dt An equal-style variable can be re-defined but only if it was originally an equal-style variable. :dd {Command input is equal-style variable} :dt Only variables that store strings can be used. :dd {Command input variable is unknown} :dt Self-explanatory. :dd {Command outputs must be specified in pairs} :dt Self-explanatory. :dd {Could not create dir for file %s\n} :dt Self-explanatory. :dd {Could not open file in print} :dt This comes from the output "print" routine of the neigh_tri command. :dd {Could not open output file %s for output object()} :dt Self-explanatory. :dd {Divide by 0 in variable formula} :dt Self-explanatory. :dd {Expected floating point parameter in variable definition} :dt The quantity being read is an integer on non-numeric value. :dd {Expected integer parameter in variable definition} :dt The quantity being read is a floating point or non-numeric value. :dd {Failed to allocate %d bytes for array %s} :dt Your OINK simulation has run out of memory. You need to run a smaller simulation or on more processors. :dd {Failed to reallocate %d bytes for array %s} :dt Your OINK simulation has run out of memory. You need to run a smaller simulation or on more processors. :dd {ID in mr command is already in use} :dt Self-explanatory. :dd {Illegal ... command} :dt Self-explanatory. Check the input script syntax and compare to the documentation for the command. You can use -echo screen as a command-line option when running OINK to see the offending line. :dd {Input line too long after variable substitution} :dt This is a hard (very large) limit defined in the input.cpp file. :dd {Input line too long: %s} :dt This is a hard (very large) limit defined in the input.cpp file. :dd {Invalid Boolean syntax in if command} :dt Self-explanatory. :dd {Invalid named command switch} :dt Only -i and -o are allowed. :dd {Invalid command-line argument} :dt One or more command-line arguments is invalid. Check the syntax of the command you are using to launch OINK. :dd {Invalid keyword in variable formula} :dt Self-explanatory. :dd {Invalid math function in variable formula} :dt Self-explanatory. :dd {Invalid seed for Marsaglia random # generator} :dt The initial seed for this random number generator must be a positive integer less than or equal to 900 million. :dd {Invalid syntax in variable formula} :dt Self-explanatory. :dd {Invalid variable evaluation in variable formula} :dt A variable used in a formula could not be evaluated. :dd {Invalid variable in next command} :dt Self-explanatory. :dd {Invalid variable name in variable formula} :dt Variable name is not recognized. :dd {Invalid variable name} :dt Variable name used in an input script line is invalid. :dd {Invalid variable style with next command} :dt Variable styles {equal} and {world} cannot be used in a next command. :dd {Label wasn't found in input script} :dt Self-explanatory. :dd {Log of zero/negative value in variable formula} :dt Self-explanatory. :dd {MR ID must be alphanumeric or underscore characters} :dt Self-explanatory. :dd {MR object map command variable is unknown} :dt Variable used as collection of strings is not recognized. :dd {MR object add comand MR object does not exist} :dt Second MR object is not recognized. :dd {MR object command input is equal-style variable} :dt Only variables that store strings can be used. :dd {MR object created by copy already exists} :dt This command creates a new MR object, which cannot already be defined. :dd {MR object map command MR object does not exist} :dt Second MR object is not recognized. :dd {Mismatch in command inputs} :dt Named command defines different number of inputs. :dd {Mismatch in command outputs} :dt Named command defines different number of outputs. :dd {Must use -in switch with multiple partitions} :dt A multi-partition simulation cannot read the input script from stdin. The -in

command-line option must be used to specify a file. :dd {Object input() invoked with invalid index} :dt Index argument must be from 1 to Ninputs. :dd {Object input() map function does not match input mode} :dt Input command is used to select map mode (mmode) which much match map() function. :dd {Object input() with no map function} :dt Input script specified input from file, but no map() method was provided by named command. :dd {Object output() as MR object not allowed} :dt Input script specied output as MR object, but named command invoked disallow flag. :dd {Object output() called for unknown MR object} :dt A MapReduce object unknown to the object manager was passed to the output() method by a named command. :dd {Object output() invoked with invalid index} :dt Index argument must be from 1 to Noutputs. :dd {Object permanent() called for unknown MR object} :dt A MapReduce object unknown to the object manager was passed to the permanent() method by a named command. :dd {Ouptut MR ID must be alphanumeric or underscore characters} :dt Self-explanatory. :dd {Power by 0 in variable formula} :dt Self-explanatory. :dd {Processor partitions are inconsistent} :dt The total number of processors in all partitions must match the number of processors LAMMPS is running on. :dd {RMAT a,b,c,d must sum to 1} :dt Self-explanatory. :dd {RMAT fraction must be < 1} :dt Self-explanatory. :dd {Sqrt of negative value in variable formula} :dt Self-explanatory. :dd {Substitution for illegal variable} :dt Input script line contained a variable that could not be substituted for. :dd {Too many edges for one vertex in reduce first_degree} :dt Number of edges of one vertex exceeds max integer in tri_find command. Will never be able to emit N^2 angles. :dd {Tour + vertex reduce exceeds one block} :dt Memory limit exceeded within sgi_enumerate command. :dd {Unbalanced quotes in input line} :dt No matching end double quote was found following a leading double quote. :dd {Universe/uloop variable count < # of partitions} :dt A universe or uloop style variable must specify a number of values >= to the number of processor partitions. :dd {Unknown command: %s} :dt The command is not known to OINK. Check the input script. :dd {Variable name must be alphanumeric or underscore characters} :dt Self-explanatory. :dd {World variable count doesn't match # of partitions} :dt A world-style variable must specify a number of values equal to the number of processor partitions. :dd Warnings: :h4,link(warn) :dlb {Placeholder} :dt No warnings are yet defined in OINK. :dd :dle "MR-MPI WWW Site"_mws - "MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line clear command :h3 [Syntax:] clear :pre [Examples:] (commands for 1st computation) clear (commands for 2nd computation) :pre [Description:] This command deletes all data structures and restores all settings to their default values. Once a clear command has been executed, it is as if OINK were starting over, with only the exceptions noted below. This command enables multiple jobs to be run sequentially from one input script. These settings are not affected by a clear command: the working directory ("shell"_shell.html command), log file status ("log"_log.html command), echo status ("echo"_echo.html command), and input script variables ("variable"_variable.html command). [Related commands:] none "MR-MPI WWW Site"_mws - "MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line named commands :h3 [Syntax:] cmomand-name params ... -i input1 input2 ... -o output1.file output1.ID output2.file output2.ID ... :pre commmand-name = name of command params = zero or more params required by command -i = start of input definitions required by command inputN = list of 0 or more input objects -o = start of output definitions to command outputN.file = list of 0 or more output files outputN.ID = list of 0 or more output MR-MPI objects :ul [Examples:] wordfreq 5 -i v_files -o NULL NULL rmat 10 8 0.25 0.25 0.25 0.25 0.0 12345 -o tmp.rmat NULL degree -i graph/edges -o degree/degree degree :pre [Description:] Invoke a named command with the list of parameters it requires, as well as the list of input and output objects it expects. In OINK a named command is a child class that derives from the Command parent class, meaning that it contains several methods that can be called from the OINK framwork. See "this section"_Section_commands.html of the manual for details on how to write new named commands and add them to OINK. The list of named commands currently included with OINK are listed on "this page"_Section_script.txt#comm. They are also listed in the source code in the file src/style_command.h which is auto-generated each time that OINK is built. Each named command has a "name", defined in the *.h file for the class, which is the command name used in the input script to invoke the command, e.g. wordfreq or rmat or degree in the examples above. Any arguments that follow the command name, upto a "-i" or "-o" argument are passed as {params} to the command before it is invoked, so that it can process and store them as needed. The number and nature of these parameters are defined by the command itself and it should

generate errors if they are not specified correctly. The code that processes parameters can be written to allow for optional parameters and keywords within the list of {params}. The "-i" and "-o" arguments can be listed in either order. The arguments that follow each of them, either between them, or upto the end of the command, are passed to an "input" and "output" processing routine within the command class. Each command requires a specific number of input and output "definitions", as explained below. Input definitions are single arguments. Output definitions are pairs of arguments. If zero input (or output) definitions are required by the command, then the "-i" (or "-o") argument need not be specified. If 2 output definitions are required, then 4 arguments must follow the "-o". Typically each required input definition is a form of data input that the command requires. It can come from reading one or more files or from an MR-MPI object that already exists. Similarly, each required output definition is a form of data output that the command produces. It can be stored either in one or more files or in an MR-MPI object that the command creates. In OINK, an MR-MPI object is a thin wrapper on a MapReduce object created via the "MR-MPI library"_md. See "this doc page"_mrmpi.html for more discussion of MR-MPI objects and input script operations that can be performed on them directly.

:line Each input definition {inputN} is one of three things. First, it can be the path name of a file or directory. Second, it can be a variable defined elsewhere in the OINK input script that contains one or more strings. In this case {inputN} should be specified as v_name, where name is the name of the variable. All the different styles of variables (except equal-style) store strings; see the "variable"_variable.html command for details. Also note that there is a "command-line option"_Section_build.txt#1_4 -var or -v which can be specified when OINK is executed to store a list of strings in an index-style variable. The strings are treated as a list of file or directory names. Thus in both the first or second case the effect is that a list of one or more file/directory names is passed to the command. The command creates a temporary, unnamed MR-MPI object and invokes a map() method within it, as specified in the code of the command class, using the list of file/directory names as input. There are several options available which influence how the list of strings specified in the input script are converted into actual file/directory paths passed to the map() method. This include wildcard characters "%" and "*". See the "input"_input.html command for details. The third option is that an input definition {inputN} can be specified as the ID of an existing MR-MPI object, which wraps a MapReduce object which contains key/value pairs. This {inputN} argument is first checked against the list of named MR-MPI objects first to determine if this is the case. If it is, then no further input processing is done. It is assumed that the data stored within the MapReduce object wrapped by the MR-MPI object is already in the form that would be produced by the map() method that would be invoked if the input definition were one or more file or directory names.

:line Each required output definition is a pair of arguments: {outputN.file} and {outputN.ID}. Either or both can be specified as NULL if no output in that form is desired. The {outputN.file} argument is the path name of a file. A map() or reduce() or scan() method, as specified in the code of the command class, will be invoked which will write data to that file when the command is finished. A processor-ID (0 to Nprocs-1) will be appended to the filename, so that when running on multiple processors, multiple files will be created. If the specified path name does not entirely exist, additional directories in the path name will be created as needed. Also, there are several options available which influence how the file name specified in the input script is converted into the file name actually opened by OINK and written to by the map(), reduce(), or scan() method. This include wildcard characters "%" and "*". See the "output"_output.html command for details. The {outputN.ID} argument is the ID of an MR-MPI object which wraps a MapReduce object. The code in the command class will have created or altered the MR-MPI object and its associated MapReduce object and populated it with data. As a final step, the specified ID is assigned to that MR-MPI object. If the ID is already in use, then the name is removed from the other MR-MPI object. When the command completes, named MR-MPI objects persist so that they can be used in subsequent input script commands. All unnamed MR-MPI objects are deleted.

:line When any named command is executed, its elapsed execution time is stored internally by OINK. This value can be accessed by the keyword "time" in an "equal-style variable"_variable.html and printed out in the following manner: variable t equal time rmat 10 8 0.25 0.25 0.25 0.25 0.0 12345 -o tmp.rmat NULL print "Time for RMat generation = \$t" :pre :line [Related commands:] "MR-MPI library commands"_mrmpi.html, "mr"_mr.html, "MR-MPI library documentation"_md, "how to write named commands"_Section_commands.html, "input"_input.html, "output"_output.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line data2graph command :h3 [Syntax:] data2graph -i input1 -o

out1.file out1.mr out2.file out2.mr out3.file out3.mr out4.file out4.mr :pre input1 = string pairs: Key = Hi Hj, Value = NULL out1 = unique vertices: Key = Hi, Value = ID (1 to N) out2 = unique edges: Key = Vi Vj, Value = NULL out3 = vertex labels: Key = Vi, Value = label out4 = edge labels: Key = Vi Vj, Value = label :ul [Examples:] data2graph -i datafiles -o vlist NULL elist NULL vatt NULL eatt NULL :pre [Description:] This is a named command which inputs raw data, typically read from files, and converts it into a labeled undirected graph, which can be further processed by other commands. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of pairs of strings, which may have many duplicates or self-edges. The input is changed by the command as it is processed further. The way that strings are read and converted into vertex and edge labels is peculiar to this command, but could be generalized. The strings are treated as hostnames extracted from network log files. The last field of the hostname is treated as the "country". If the country has 3 characters is is considered to be the US, else it is a foreign country. As the raw data is read, US/US edges and edges with hostnames that have only one field are discarded. A vertex is labeled with a 1 if it is US, a 2 if it is foreign. An edge is labeled with a 1 if its 2 vertices are the same country, else with a 2. Out1 will store the list of unique vertices, with the original string and a vertex ID (Vi = 1 to N) assigned to each. Out2 will store the list of unique edges, with no duplicates or self-edges, as pairs of vertex IDs. Out3 will store the list of vertex labels, where label is an integer. Out4 will store the list of edge labels, where label is an integer. [Related commands:] "data2rare"_data2rare.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line data2rare command :h3 [Syntax:] data2rare Nthresh -i input1 -o out1.file NULL out2.file NULL :pre Nthresh = define "rare" as visits <= Nthresh input1 = dated domain/user pairs: Key = domain data, Value = user out1 = per-domain statistics written to out1.file (see below) out2 = per-user statistics written to out2.file (see below) :ul [Examples:] data2rare 1 -i datafiles -o domain.file NULL user.file NULL :pre [Description:] This is a named command which inputs raw data, typically read from files. Individual records have 3 fields: domain, user, data. A record represents a user visiting a domain on a date. This data is processed into 2 statistical forms, as described below, and output to files. The user statistics include a count of how many "rare" domains the user visited. Rare is defined as a domain visited (on a particular date) by {Nthresh} or less unique users. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of 3 strings, which may have many duplicates or self-edges. The input is changed by the command as it is processed further. The way that the 3 strings are read and processed into statistics is peculiar to this command, but could be generalized. The 3 strings are a domain, user, and date. What is read from the file is actually a full hostname. A subset is treated as the domain. See the code for details of how the domain is extracted. Statistics are generated and output in the following format: The per-domain statistics file has one line per unique domain per date with these fields: date,domain,n1,n2 n1 = total visits to domain on date n2 = unique users who visited domain on date :pre N2 should be <= N1 for each entry. The per-user statistics file has one line per unique user per date with these fields: date,user,n1,n2,n3 n1 = total visits by user on date n2 = unique domains visited by user on date n3 = rare domains visited by user on date :pre N3 should be <= N2 and N2 <= N1 for each entry. Out1 will store the per-domain statistics. Out2 will store the per-user statistics. Output of this data to MR objects is not allowed. Only files can be output by this command. [Related commands:] "data2graph"_data2graph.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line degree command :h3 [Syntax:] degree -i input1 -o out1.file out1.mr :pre input1 = graph edges: Key = Vi Vj, Value = NULL out1 = degree of each vertex: Key = Vi, Value = degree :ul [Examples:] degree -i mrv -o degree.list NULL [Description:] This is a named command which calculates the degree of each vertex in an undirected graph. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of edges, assumed to have no duplicates or self-edges. The input is unchanged by this command. A count of the number of edges containing each vertex is calculated. Out1 will store the degree count of each vertex. [Related commands:] "neighbor"_neighbor.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line echo

command :h3 [Syntax:] echo style :pre style = {none} or {screen} or {log} or {both} :ul [Examples:] echo both echo log :pre [Description:] This command determines whether OINK echoes each input script command to the screen and/or log file as it is read and processed. If an input script has errors, it can be useful to look at echoed output to see the last command processed. The "command-line switch" _Section_start.html#2_4 -echo can be used in place of this command. The default is echo log, i.e. commands are echoed to the log file. [Related commands:] none "MR-MPI WWW Site" _mws - "MR-MPI Documentation" _md - "OINK Documentation" _od - "OINK Commands" _oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,../doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line if command :h3 [Syntax:] if boolean then t1 t2 ... elif boolean f1 f2 ... elif boolean f1 f2 ... else e1 e2 ... :pre boolean = a Boolean expression evaluated as TRUE or FALSE (see below) then = required word t1,t2,...,tN = one or more OINK commands to execute if condition is met, each enclosed in quotes elif = optional word, can appear multiple times f1,f2,...,fN = one or more OINK commands to execute if elif condition is met, each enclosed in quotes (optional arguments) else = optional argument e1,e2,...,eN = one or more OINK commands to execute if no condition is met, each enclosed in quotes (optional arguments) :ul [Examples:] if "\$\{steps\} > 1000" then exit if "\$x <= \$y" then "print X is smaller = \$x" else "print Y is smaller = \$y" if "(\$\{flag\} == 0) || (\$n < 1000)" then & "graph reduce myfunc" & elif "\$\{flag\} == 1" & "graph reduce myfunc2" & else & "graph kmv_stats 2" & "print Elapsed time = \$t" if "\$\{niter\} > \$\{niter_previous\}" then "jump file1" else "jump file2" :pre [Description:] This command provides an in-then-else capability within an input script. A Boolean expression is evaluated and the result is TRUE or FALSE. Note that as in the examples above, the expression can contain variables, as defined by the "variable" _variable.html command, which will be evaluated as part of the expression. Thus a user-defined formula that reflects the current state of the simulation can be used to issue one or more new commands. If the result of the Boolean expression is TRUE, then one or more commands (t1, t2, ..., tN) are executed. If it is FALSE, then Boolean expressions associated with successive elif keywords are evaluated until one is found to be true, in which case its commands (f1, f2, ..., fN) are executed. If no Boolean expression is TRUE, then the commands associated with the else keyword, namely (e1, e2, ..., eN), are executed. The elif and else keywords and their associated commands are optional. If they aren't specified and the initial Boolean expression is FALSE, then no commands are executed. The syntax for Boolean expressions is described below. Each command (t1, f1, e1, etc) can be any valid OINK input script command. If the command is more than one word, it must be enclosed in quotes, so it will be treated as a single argument, as in the examples above. IMPORTANT NOTE: If a command itself requires a quoted argument (e.g. a "print" _print.html command), then double and single quotes can be used and nested in the usual manner, as in the examples above and below. See "this section" _Section_commands.html#3_2 of the manual for more details on using quotes in arguments. Only one level of nesting is allowed, but that should be sufficient for most use cases. Note that by using the line continuation character the if command can be spread across many lines, though it is still a single command: if "\$a < \$b" then & "print 'Minimum value = \$a'" & "graph ..." & else & 'print "Minimum value = \$b"' & "graph ..." :pre Note that if one of the commands to execute is an invalid OINK command, such as "exit" in the first example above, then executing the command will cause OINK to halt. Note that by jumping to a label in the same input script, the if command can be used to break out of a loop. See the "variable delete" _variable.html command for info on how to delete the associated loop variable, so that it can be re-used later in the input script. Here is an example of a double loop which uses the if and "jump" _jump.html commands to break out of the inner loop when a condition is met, then continues iterating thru the outer loop. label loopa variable a loop 5 label loopb variable b loop 5 print "A,B = \$a,\$b" ... if "\$b > 2" then "print 'Jumping to another script'" "jump in.script break" next b jump in.script loopb label break variable b delete :pre next a jump in.script loopa :pre :line The Boolean expressions for the if and elif keywords have a C-like syntax. Note that each expression is a single argument within the if command. Thus if you want to include spaces in the expression for clarity, you must enclose the entire expression in quotes. An expression is built out of numbers: 0.2, 100, 1.0e20, -15.4, etc :pre and Boolean operators: A == B, A != B, A < B, A <= B, A > B, A >= B, A & B, A || B, !A :pre Each A and B is a number or a variable reference like \$a or \$\{abc\}, or another Boolean expression. If a variable is used it must produce a number when evaluated and substituted for in the expression, else an error will be generated. Expressions are evaluated left to right and have the usual C-style precedence: the unary logical NOT operator "!" has the highest precedence, the 4 relational operators "==", "<", "<=", ">", ">=" are next; the two remaining relational operators "!=" and "!=" are next; then the logical AND operator "&" and finally the logical OR operator "||" has the lowest precedence. Parenthesis can be used to

group one or more portions of an expression and/or enforce a different order of evaluation than what would occur with the default precedence. The 6 relational operators return either a 1.0 or 0.0 depending on whether the relationship between x and y is TRUE or FALSE. The logical AND operator will return 1.0 if both its arguments are non-zero, else it returns 0.0. The logical OR operator will return 1.0 if either of its arguments is non-zero, else it returns 0.0. The logical NOT operator returns 1.0 if its argument is 0.0, else it returns 0.0. The overall Boolean expression produces a TRUE result if the result is non-zero. If the result is zero, the expression result is FALSE. [Related commands:] "variable"_variable.html, "print"_print.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c

:link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line include command :h3 [Syntax:] include file :pre file = filename of new input script to switch to :ul [Examples:] include newfile include in.run2 :pre [Description:] This command opens a new input script file and begins reading OINK commands from that file. When the new file is finished, the original file is returned to. Include files can be nested as deeply as desired. If input script A includes script B, and B includes A, then OINK could run for a long time. If the filename is a variable (see the "variable"_variable.html command), different processor partitions can run different input scripts. [Related commands:] "variable"_variable.html, "jump"_jump.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c

:link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line input command :h3 [Syntax:] input N keyword value ... :pre N = which input to set options for :ulb,l one or more keyword/value pairs may be appended :l keyword = {prepend} or {substitute} or {multi} or {mmode} or {recurse} or {self} or {readfile} or {nmap} or {sepchar} or {sepstr} or {delta} {prepend} value = string to prepend to file/directory path names {substitute} value = 0 or 1 = how to substitute for "%" in path name {multi} value = Nmulti = multiplicity of path names to generate {mmode} value = 0 or 1 or 2 = which style of map() method to invoke {recurse} value = 0 or 1 = passed to map() method {self} value = 0 or 1 = passed to map() method {readfile} value = 0 or 1 = passed to map() method {nmap} value = number of map tasks = passed to map() method {sepchar} value = single character = passed to map() method {sepstr} value = string = passed to map() method {delta} value = Ndelta = passed to map() method :pre :ule [Examples:] input 1 multi 4 input 2 self 1 substitute 4 prepend /scratch%/hadoop-datastore/local_files :pre [Description:] This command is used to control the reading of input data that a "named command"_command.html performs as part of its input. It does this by setting options on specific inputs to "named commands"_command.html. The options set by this command are in effect for ONLY the next named command. After a named command is invoked, it restores all input options to their default values. Note that all of the options which can be set by this command have default values, so you don't need to set those you don't want to change. As described on the "named command"_command.html doc page, a named command may specify one or more input descriptors. If the descriptor is one or more file or directory names, then each of them is converted into a list of strings which is passed to a map() method of a created MR-MPI object, along with various other arguments needed by the map() method. The purpose of the input command is to give you control over how that conversion takes place and what the values of those additional arguments are. The {N} value corresponds to a particular input descriptor, as defined by the "named command"_command.html. It should be an integer from 1 to Ninput, where Ninput is the number of input descriptors the command requires. The input command can be used multiple times with the same {N} to specify different parameters, e.g. one at a time. The remaining arguments are pairs of {keywords} and {values}. One or more can be specified. :line The {prepend}, {substitute}, and {multi} keywords alter the file and directory path names specified with an input descriptor in a named command.

IMPORTANT NOTE: The {prepend}, {substitute}, and {multi} keywords are applied to each file of directory name in the list of such names that the "named command"_command.html uses in its input descriptor.

IMPORTANT NOTE: The {prepend} and {substitute} keywords can also be set globally so that their values will be applied to all input descriptors of all "named commands"_command.html. See the "set"_set.html command for details. If an input command is not used to override the global value, then the global value is used by the "named command"_command.html. The {prepend} keyword specifies a path name to prepend to each input file or directory name specified with the "named command"_command.html. The prepend string is presumed to be a directory name and should be specified without the trailing "/" character, since that is added when the prepending is done. Input file or directory names specified with the "named command"_command.html can contain either or

both of two wildcard characters: "%" or "*". Only the first occurrence of each wildcard character is replaced. If the {substitute} keyword is set to 0, then a "%" is replaced by the processor ID, 0 to Nprocs-1. If it is set to N > 0, then "%" is replaced by (proc-ID % N) + 1. I.e. for 8 processors and N = 4, then the 8 processors replace the "%" with (1,2,3,4,1,2,3,4). This can be useful for multi-core nodes where each core has its own local disk. E.g. you wish each core to read data from one disk. If a "*" appears, then the file or directory name is duplicated N times where N is the value set by the {multi} keyword. In each of the N duplicates, the "*" is replaced by the number 1 to N. Again, this can be useful for multi-core nodes where each core has its own local disk. E.g. you want a single core to read data from each of several local disks on the node, presumably because you have launched an MPI job so that it runs on a single core per node.

The {mmode} keyword stands for "map mode" and determines what form of the "MR-MPI library map() method" [../doc/map.html](http://www.sandia.gov/~sjplimp/mapreduce.html) is invoked by the "named command" [_command.html](#). It is up to the coding of the "named command" to determine which of these forms of data input it supports. There are 3 variants of the map() method which involve file input: mmode = 0 = map(int nstr, char **strings, int self, int recurse, int readfile, void (*mymap)(), void *ptr) mmode = 1 = map(int nmap, int nstr, char **strings, int recurse, int readfile, char sepchar, int delta, void (*mymap)(), void *ptr) mmode = 2 = map(int nmap, int nstr, char **strings, int recurse, int readfile, char *sepstr, int delta, void (*mymap)(), void *ptr)

The "nstr" and "strings" arguments to these methods are created by OINK, using the settings described above. The remaining arguments are set by the keywords of the input command, as needed. Note that some keywords have no meaning for certain map() method variants, in which case they are simply ignored. The meaning of the {self}, {recurse}, {readfile}, {nmap}, {sepchar}, {sepstr}, and {delta} keywords is the same as explained on the doc page for the "map() method" [../doc/map.html](http://www.sandia.gov/~sjplimp/mapreduce.html) of the MR-MPI library. The value for the {sepchar} keyword will be treated as a single character. The value for the {sepstr} keyword will be treated as a string.

[Related commands:] "output" [_output.html](#), "named commands" [_command.html](#), "how to write named commands" [_Section_commands.html](#), "set" [_set.html](#) [Defaults:] The option defaults are prepend = NULL, substitute = 0, multi = 1, mmode = 0, recurse = 0, self = 0, readfile = 0, nmap = 0, sepchar = newline character, sepstr = newline, delta = 80. "MR-MPI WWW Site" [_mws](#) - "MR-MPI Documentation" [_md](#) - "OINK Documentation" [_od](#) - "OINK Commands" [_oc](#) :c :link(mws,<http://www.sandia.gov/~sjplimp/mapreduce.html>) :link(md,[../doc/Manual.html](#)) :link(od,[Manual.html](#)) :link(oc,[Section_commands.html#comm](#))

jump command :h3 [Syntax:] jump file label :pre file = filename of new input script to switch to label = optional label within file to jump to :ul [Examples:] jump newfile jump in.run2 runloop jump SELF runloop :pre [Description:] This command closes the current input script file, opens the file with the specified name, and begins reading OINK commands from that file. Unlike the "include" [_include.html](#) command, the original file is not returned to, although by using multiple jump commands it is possible to chain from file to file or back to the original file. If the word "SELF" is used for the filename, then the current input script is re-opened and read again.

IMPORTANT NOTE: The SELF option is not guaranteed to work when the current input script is being read through stdin (standard input), e.g. `lmp_g++ < in.script` :pre since the SELF option invokes the C-library rewind() call, which may not be supported for stdin on some systems. This can be worked around by using the "-in command-line argument" [_Section_start.html#2_6](#) or the "-var command-line argument" [_Section_start.html#2_6](#) to pass the script name as a variable to the input script In the latter case, the "fname" "variable" [_variable.html](#) could be used in place of SELF. E.g. `lmp_g++ -in in.script :pre lmp_g++ -var fname n.script < in.script :pre` The 2nd argument to the jump command is optional. If specified, it is treated as a label and the new file is scanned (without executing commands) until the label is found, and commands are executed from that point forward. This can be used to loop over a portion of the input script, as in this example. These commands perform 10 runs, each of 10000 steps, and create 10 dump files named file.1, file.2, etc. The "next" [_next.html](#) command is used to exit the loop after 10 iterations. When the "a" variable has been incremented for the tenth time, it will cause the next jump command to be skipped.

```
variable a loop 10 label loop dump 1 all atom 100 file.$a run 10000 undump 1 next
a jump in.lj loop :pre
```

If the jump {file} argument is a variable, the jump command can be used to cause different processor partitions to run different input scripts. In this example, LAMMPS is run on 40 processors, with 4 partitions of 10 procs each. An in.file containing the example variable and jump command will cause each partition to run a different simulation.

```
mpirun -np 40 lmp_ibm -partition 4x10 -in in.file :pre variable f world
script.1 script.2 script.3 script.4 jump $f :pre
```

Here is an example of a double loop which uses the "if" [_if.html](#) and jump commands to break out of the inner loop when a condition is met, then continues iterating thru the outer loop.

```
label loopa variable a loop 5 label loopb variable b loop 5 print "A,B = $a,$b" run 10000 if $b > 2 then
```

"jump in.script break" next b jump in.script loopb label break variable b delete :pre next a jump in.script loopa :pre IMPORTANT NOTE: If you jump to a file and it does not contain the specified label, OINK will come to the end of the file and exit. [Related commands:] "variable"_variable.html, "include"_include.html, "label"_label.html, "next"_next.html "MR-MPI WWW Site"_mws – "MR-MPI Documentation"_md – "OINK Documentation"_od – "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line label command :h3 [Syntax:] label ID :pre ID = string used as label name :ul [Examples:] label xyz label loop :pre [Description:] Label this line of the input script with the chosen ID. Unless a jump command was used previously, this does nothing. But if a "jump"_jump.html command was used with a label argument to begin invoking this script file, then all command lines in the script prior to this line will be ignored. I.e. execution of the script will begin at this line. This is useful for looping over a section of the input script as discussed in the "jump"_jump.html command. [Related commands:] none "MR-MPI WWW Site"_mws – "MR-MPI Documentation"_md – "OINK Documentation"_od – "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line log command :h3 [Syntax:] log file :pre file = name of new logfile :ul [Examples:] log log.graph :pre [Description:] This command closes the current OINK log file, opens a new file with the specified name, and begins logging information to it. If the specified file name is {none}, then no new log file is opened. If multiple processor partitions are being used, the file name should be a variable, so that different processors do not attempt to write to the same log file. The file "log.oink" is the default log file for a OINK run. The name of the initial log file can also be set by the command-line switch –log. See "this section"_Section_start.html#2_6 for details. The default OINK log file is named log.oink [Related commands:] none "MR-MPI WWW Site"_mws – "MR-MPI Documentation"_md – "OINK Documentation"_od – "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line mr command :h3 [Syntax:] mr MR-ID verbosity timer memsize outofcore :pre MR-ID = ID of new MR-MPI object to create verbosity = verbosity setting (optional) timer = timer setting (optional) memsize = memsize setting (optional) outofcore = outofcore setting (optional) :ul [Examples:] mr edge mr edge 2 mr edge 1 1 16 0 :pre [Description:] Create a new MR-MPI object, which can be referenced by name elsewhere in your input script. In OINK, a MR-MPI object is simply a thin wrapper on a MapReduce object created via the "MR-MPI library"_md. The MR-MPI object has an ID which can be used elsewhere in the input script. For example it can be used as input to a "named command"_command.html or in a "MR-MPI library command"_mrmapi.html. The ID of an MR-MPI object can only contain alphanumeric characters and underscores. When the underlying MapReduce object is created, it will have default settings as described "here"../doc/settings.html. Several of these settings can be overridden by the 4 options listed above. If none of them are specified, then the default settings are used. To reset one of the settings, you must specify all the settings that precede it. E.g. if just two optional arguments are used, they are the verbosity and timer settings. [Related commands:] "MR-MPI library commands"_mrmapi.html, "named commands"_command.html "MR-MPI WWW Site"_mws – "MR-MPI Documentation"_md – "OINK Documentation"_od – "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line MR-MPI library commands :h3 [Syntax:] MR-ID keyword args ... :pre MR-ID = ID of previously created MR object keyword = MR-MPI library function to invoke args = arguments to library function :ul [Examples:] edge map/task 100 mymap edge map/task 100 mymap 1 edge collate NULL edge reduce myreduce edge kv_stats 1 edge set timer 1 [Description:] Invoke a MR-MPI library functions directly on a previously created MR-MPI objects. In OINK, an MR-MPI object is a thin wrapper on a MapReduce object created via the "MR-MPI library"_md. They can be created by the "mr"_mr.html command or can be output by a "named command"_command.html. Such an MR-MPI object has an ID which is the command name used in the input script to trigger the library calls, e.g. "edge" in the examples above. The keyword is the library function to invoke on the underlying MapReduce object wrapped by the MR-MPI object. These have a one-to-one correspondence with the methods available in the "MR-MPI library"_md. Here is the list of keywords and their arguments. The arguments used in the OINK input script correspond to the arguments used by each library method. Arguments in parentheses are optional. More details are discussed below. Keywords, Arguments, delete, none, copy, MR2-ID, add, MR2-ID, aggregate, NULL or hash-function, broadcast, root, clone, none, close, none, collapse, type key, collate, NULL of hash-function,

compress, reduce-function, convert, none, gather, nprocs, map/task, nmap map-function (addflag), map/char, nmap strings recurse readfile sepchar delta map-function (addflag), map/string, nmap strings recurse readfile sepstr delta map-function (addflag), map/mr, MR2-ID map-function (addflag), open, none, print, (file) (fflag) proc nstride kflag vflag, reduce, reduce-function, scan/kv, scan-function, scan/kmv, scan-function, scrunch, nprocs type key, sort_keys, flag or compare-function, sort_values, flag or compare-function, sort_multivalues, flag or compare-function, kv_stats, level, kmv_stats, level, cumulative_stats, level reset, set, name value :tb(c=2) The MR2-ID used as an argument to the "copy", "add", and "map/mr" keywords should be the ID of another previously defined MR-MPI object. IMPORTANT NOTE: The syntax for the copy keyword in an OINK script is as follows: MR-ID copy MR2-ID. This creates a new MR-MPI object MR2-ID, which is a copy of the existing MR-MPI object MR-ID. The MR2-ID object cannot already exist. This corresponds to the following C++ calling syntax for the "copy()" method of the MR-MPI library "[../doc/copy.html](#)", but note that the OINK syntax is somewhat reversed: `MapReduce *mr2 = mr->copy();` ;pre The map-function, reduce-function, hash-function, compare-function, scan-function arguments to various keywords are the names of functions that will be called back to by the MR-MPI library. Within OINK, these must be names of functions defined in `map_*.cpp`, `reduce_*.cpp`, `hash_*.cpp`, `compare_*.cpp`, or `scan_*.cpp` files with the appropriate function prototype. When you build OINK, these files are scanned, the function prototypes extracted, and the `style_map.h`, `style_reduce.h`, `style_hash.h`, `style_compare.h`, `style_scan.h` files are created which enables a function name you list in your input script to be recognized by OINK. Note that as new `map()`, `reduce()`, etc functions are added to the OINK src directory, they automatically become available to your script to use in MR-MPI library commands. Thus you can use OINK to accumulate a collection of useful `map()`, `reduce()`, etc functions. These functions can also be used with "named commands" `_command.html` as discussed "here" `_Section_command.html`. Note that `map()` functions come in 4 different flavors, with different prototypes, as "detailed here" "[../doc/map.html](#)". Which you should use depends on which map variant you invoke, i.e. `map/task`, `map/char`, `map/string`, or `map/mr`. Likewise, `scan()` functions come in 2 different flavors, as "detailed here" "[../doc/scan.html](#)", one for use with `scan/kv` and the other with `scan/kmv`. The "strings" argument to the `{map/char}` and `{map/string}` keywords can take one of two forms. It can be a single filename or directory. If the latter, then the "map()" method in the MR-MPI library "[../doc/map.html](#)" reads the files in the directory. Or it can be a variable defined elsewhere in the OINK input script that contains one or more strings which are passed to the `map()` method as a collection of strings. In this case the "strings" argument should be specified as `v_name`, where name is the name of the variable. All the different styles of variables (except equal-style) store strings; see the "variable" `_variable.html` command for details. Also note that there is a "command-line option" `_Section_build.txt#1_4 -var` or `-v` which can be specified when OINK is executed to store a list of filenames in an index-style variable. The `sepchar` and `sepstr` arguments to the `{map/char}` and `{map/string}` keywords should be a single character or a string of characters. The `addflag` argument to the various `{map}` keywords is optional. It should be 1 if you wish to add key/value pairs to those already contained in a MapReduce object. The type argument to the `{collapse}` and `{scrunch}` keywords should be one of the following: "int", "uint46", "double", or "str". The key that follows will be converted into that data type to use as the key argument to the MR-MPI library function. The `{print}` keyword takes either 4 or 6 arguments. If 6 are used, the first two are a file name and file flag, the same as is available with the "print()" method in the MR-MPI library "[../doc/print.html](#)". The flag argument to the various `{sort}` keywords is an integer (e.g. 1 or -1) that can be used in place of a compare-function. This is the same integer that the "sort methods in the MR-MPI library" "[../doc/sort.html](#)" takes as a valid argument. The `{set}` keyword takes a "name" and "value" argument. These can be any of the options that are valid to set for a MapReduce object in the MR-MPI library, as "discussed here" "[../doc/settings.html](#)". E.g. the command "edge verbosity 1" will set the verbosity level to 1 in the MapReduce object wrapped by the MR-MPI object named "edge". IMPORTANT NOTE: There is currently no way in OINK to pass a data pointer to the various MR-MPI library functions that accept it, e.g. to `map()` or `reduce()`. When using the library from a programming language, such as C++ or C, this is powerful option for passing extra information to the user callback `map()` or `reduce()` function. We are still thinking about the best way to do this, at least in some limited fashion, from an OINK input script. When any MR-MPI library command is executed, its elapsed execution time is stored internally by OINK. This value can be accessed by the keyword "time" in an "equal-style variable" `_variable.html` and printed out in the following manner: variable t equal time edge map/task 100 mymap print "Time for map/task = \$t" ;pre :line [Related commands:] "named commands" `_command.html`, "mr" `_mr.html`, "MR-MPI library documentation" `_md`, "map()", `reduce()`, etc

functions"_Section_functions.txt "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..../doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line neigh_tri command :h3 [Syntax:] neigh_tri dirname -i input1 input2 -o out1.file out1.mr :pre dirname = directory name to create set of output files in, one per vertex input1 = graph neighbors: Key = Vi, Value = Vj Vk ... input2 = triangles: Key = Vi Vj Vk, Value = NULL out1 = neighbors + triangle edges of each vertex: Key = Vi, MultiValue = Vj Vk ... (Vj Vk) (Vm Vn) ... :ul [Examples:] neigh_tri myneigh -i mrn mrtri -o NULL mrnplus [Description:] This is a named command which calculates a list of edges associated with each vertex in an undirected graph, which include all edges the vertex is in (its first neighbors) and also edges between pairs of its first neighbors (triangle edges). This set of data is written to a file per vertex as a list of edges. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of neighbors of each vertex. See the "neighbor"_neighbor.html command, which can compute this data. Input2 stores a set of triangles. See the "tri_find"_tri_find.html command, which can compute this data. The two inputs are unchanged by this command. These 2 data sets are merged to identify the edges that exist between pairs of neighbors of each vertex. This information is written to a file per vertex. The name of each file is dirname/Vi where {dirname} is the specified argument (a directory name), and Vi is the vertex ID. Each file will contain a list of edges, one per line, written as Vm Vn. For some of the Vm will equal Vi, which means they are edges containing Vi, i.e. they are the first neighbors of Vi. Other edges will have Vm != Vi. These are edges between pairs of first neighbors. Out1 will store the neighbor and triangle edge information as key/multivalue (KMV) pairs, not as key/value (KV) pairs (the usual form of output). Out1.file must be specified as NULL with the "-o" argument so that the output is only allowed as an MR-MPI object, not as a file. This is because the file would contain data for all the vertices together. The equivalent info is already output as one file per vertex, as described above. NOTE: alter the neigh_tri.cpp code so that it uses the input dirname with expandpath() to apply the global prepend and substitute settings ?? [Related commands:] "neighbor"_neighbor.html, "tri_find"_tri_find.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..../doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line neighbor command :h3 [Syntax:] neighbor -i input1 -o out1.file out1.mr :pre input1 = graph edges: Key = Vi Vj, Value = NULL out1 = neighbors of each vertex: Key = Vi, Value = Vj Vk ... :ul [Examples:] degree -i mrv -o degree.list NULL [Description:] This is a named command which calculates the list of neighbors of each vertex in an undirected graph. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of edges, assumed to have no duplicates or self-edges. The input is unchanged by this command. A list of all the vertices each vertex shares an edge with is calculated. These are the first neighbors of each vertex. Out1 will store the list of neighbor vertices for each vertex. Thi list is a single value in a key/value pair which is a vector of vertex IDs, one after the other. [Related commands:] "degree"_degree.html, "neigh_tri"_neigh_tri.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..../doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line next command :h3 [Syntax:] next variables :pre variables = one or more variable names :ul [Examples:] next x next a t x myTemp :pre [Description:] This command is used with variables defined by the "variable"_variable.html command. It assigns the next value to the variable from the list of values defined for that variable by the "variable"_variable.html command. Thus when that variable is subsequently substituted for in an input script command, the new value is used. See the "variable"_variable.html command for info on how to define and use different kinds of variables in OINK input scripts. If a variable name is a single lower-case character from "a" to "z", it can be used in an input script command as \$a or \$z. If it is multiple letters, it can be used as \${myTemp}. If multiple variables are used as arguments to the {next} command, then all must be of the same variable style: {index}, {loop}, {universe}, or {uloop}. An exception is that {universe}- and {uloop}-style variables can be mixed in the same {next} command. All the variables specified with the next command are incremented by one value from their respective list or values. {String-} or {equal}- or {world}-style variables cannot be used with the the next command, since they only store a single value. When any of the variables in the next command has no more values, a flag is set that causes the input script

to skip the next "jump" `_jump.html` command encountered. This enables a loop containing a next command to exit. As explained in the "variable" `_variable.html` command, the variable that has exhausted its values is also deleted. This allows it to be used and re-defined later in the input script. When the next command is used with {index}– or {loop}–style variables, the next value is assigned to the variable for all processors. When the next command is used with {universe}– or {uloop}–style variables, the next value is assigned to whichever processor partition executes the command first. All processors in the partition are assigned the same value. Running OINK on multiple partitions of processors via the "-partition" command-line switch is described in "this section" `_Section_build.html#1_4` of the manual. {Universe}– and {uloop}–style variables are incremented using the files "tmp.oink.variable" and "tmp.oink.variable.lock" which you will see in your directory during such a OINK run. Here is an example of running a series of simulations using the next command with an {index}–style variable. If this input script is named `in.graph`, 8 simulations would be run using data files from directories `run1` thru `run8`.

```
variable d index run1 run2 run3 run4 run5 run6 run7 run8 shell cd $d graph -i data.graph shell cd ..
clear next d jump in.graph :pre If the variable "d" were of style {universe}, and the same in.graph input script
were run on 3 partitions of processors, then the first 3 simulations would begin, one on each set of processors.
Whichever partition finished first, it would assign variable "d" the 4th value and run another simulation, and so
forth until all 8 simulations were finished. Jump and next commands can also be nested to enable multi-level
loops. For example, this script will run 15 simulations in a double loop. variable i loop 3 variable j loop 5 clear ...
print Running simulation $i.$j graph -i data.polymer.$i$j next j jump in.script next i jump in.script :pre Here is an
example of a double loop which uses the "if" _if.html and "jump" _jump.html commands to break out of the inner
loop when a condition is met, then continues iterating thru the outer loop. label loopa variable a loop 5 label loopb
variable b loop 5 print "A,B = $a,$b" run 10000 if $b > 2 then "jump in.script break" next b jump in.script loopb
label break variable b delete :pre next a jump in.script loopa :pre [Related commands:] "jump" _jump.html,
"include" _include.html, "shell" _shell.html, "variable" _variable.html, "MR-MPI WWW Site" _mws – "MR-MPI
Documentation" _md – "OINK Documentation" _od – "OINK Commands" _oc :c
:link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html)
:link(oc,Section_commands.html#comm) :line output command :h3 [Syntax:] output N keyword value ... :pre N =
which output to set options for :ulb,l one or more keyword/value pairs may be appended :l keyword = {prepend}
or {substitute} {prepend} value = string to prepend to file path names {substitute} value = 0 or 1 = how to
substitute for "%" in path name :ule [Examples:] output 1 substitue 4 output 2 substitute 4 prepend
/scratch%/hadoop-datastore/local_files :pre [Description:] This command is used to control the writing of data
that a "named command" _command.html performs as part of its output. It does this by setting options on specific
outputs to "named commands" _command.html. The options set by this command are in effect for ONLY the next
named commmand. After a named command is invoked, it restores all output options to their default values. Note
that all of the options which can be set by this command have default values, so you don't need to set those you
don't want to change. As described on the "named command" _command.html doc page, a named command may
specify one or more output descriptors. Each descriptor is a pair of arguemnts, the first of which is an output
filename (if it is not specified as NULL). OINK converts the specified argument into an actual filename which is
opened by each processor. The purpose of the output command is to give you control over how that conversion
takes place. The {N} value corresponds to a particular output descriptor, as defined by the "named
command" _command.html. It should be an integer from 1 to Noutput, where Noutput is the number of output
descriptors the command requires. The output command can be used multiple times with the same {N} to specify
different parameters, e.g. one at a time. The remaining arguments are pairs of {keywords} and {values}. One or
more can be specified. :line The {prepend} and {substitute} keywords alter the file and directory path names
specified with the filename of an output descriptor in a named command. IMPORTANT NOTE: The {prepend}
and {substitute} keywords can also be set globally so that their values will be applied to all output descriptors of
all "named commands" _command.html. See the "set" _set.html command for details. If an output command is not
used to override the global value, then the global value is used by the "named command" _command.html. The
{prepend} keyword specifies a path name to prepend to the output file specified with the "named
command" _command.html. The prepend string is presumed to be a directory name and should be specified
without the trailing "/" character, since that is added when the prepending is done. Ouptut file or directory names
specified with the "named command" _command.html can contain either or both of two wildcard characters: "%"
or "*". Only the first occurrence of each wildcard character is replaced. If the {substitute} keyword is set to 0,

```

then a "%" is replaced by the processor ID, 0 to Nprocs-1. If it is set to N > 0, then "%" is replaced by (proc-ID % N) + 1. I.e. for 8 processors and N = 4, then the 8 processors replace the "%" with (1,2,3,4,1,2,3,4). This can be useful for multi-core nodes where each core has its own local disk. E.g. you wish each core to write data to one disk. IMPORTANT NOTE: The processor ID is also added as a suffix to the specified output file by each processor, so that one output file per processor is generated. This is in addition to any replacement of a "%" wildcard character. If a "*" appears, then it is replaced with a 1. Unlike for "input files" _input.html, this is not a particularly useful wildcard for output files. :line [Related commands:] "input" _input.html, "named commands" _command.html, "how to write named commands" _Section_commands.html, "set" _set.html [Defaults:] The option defaults are prepend = NULL, substitute = 0, multi = 1, mmode = 0, recurse = 0, self = 0, readfile = 0, nmap = 0, sepchar = newline character, sepstr = newline, delta = 80. "MR-MPI WWW Site" _mws - "MR-MPI Documentation" _md - "OINK Documentation" _od - "OINK Commands" _oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line print command :h3 [Syntax:] print str :pre str = text string to print, which may contain variables :ul [Examples:] print "Done with first stage" print "Elapsed time = \$t secs on \$p procs" [Description:] Print a text string to the screen and logfile. One line of output is generated. If the string has white space in it (spaces, tabs, etc), then you must enclose it in quotes so that it is treated as a single argument. If variables are included in the string, they will be evaluated and their current values printed. See the "variable" _variable.html command for a description of various kinds of variables, any of which can be used with the print command. Note that there are keywords for the number of processors and elapsed time for a command or MR-MPI library call which can be accessed with variables, e.g. variable t equal time variable p equal nprocs print "Elapsed time = \$t secs on \$p procs" :pre [Related commands:] "variable" _variable.html "MR-MPI WWW Site" _mws - "MR-MPI Documentation" _md - "OINK Documentation" _od - "OINK Commands" _oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line rmat command :h3 [Syntax:] rmat N Nz a b c d fraction seed -o out1.file out1.mr :pre N = order of matrix, 2^N = number of rows in matrix Nz = average # of non-zeroes per row, Nz * 2^N = total # of non-zeroes a,b,c,d = RMAT parameters which sum to 1.0 fraction = RMAT twiddle factor seed = random number seed (positive integer) out1 = graph edges: Key = Vi Vj, Value = NULL :ul [Examples:] rmat 20 8 0.45 0.25 0.25 0.05 0.0 284958 -o NULL mre [Description:] Generate a sparse random matrix via the rules defined for RMAT matrices, as discussed in the paper by "(Christoff)" _#Christoff. Such matrices are often used to represent graphs where the vertices are numbered 1 to Nrows, and the non-zero matrix entries represent edges. The number of rows and non-zero entries are determined by the specified {N} and {Nz} arguments. Depending on the choice of the RMAT parameters the degree distribution of the resulting graph can be roughly uniform or highly skewed, which is useful in modeling different kind of graphs, e.g. Internet connectivity. The a,b,c,d parameters must sum to 1.0 and represent weighting for the 4 different quadrants of the matrix. As non-zero entries are generated, they are assigned to each quadrant in a recursive manner using the a,b,c,d weightings and a random number generator. A fraction value of 0.0 means the a,b,c,d weightings are used as-is. A fraction value > 0.0 but < 1.0 means the weightings are randomly twiddled at each iteration of the recursion. The MapReduce algorithm for performing the RMAT generation is described in the paper by "(Plimpton)" _#Plimpton. This rmat command implements the first of the two algorithms in the Plimpton and Devine paper. See the "named command" _command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. This command takes no inputs via the "-i" argument for "named commands" _command.html. Out1 will store the list of edges of the resultig RMAT graph, or equivalently, the I,J indices of non-zeroes in the matrix. There will be exactly Nz * 2^N entries in out1. This may include some self-edges (I,I) and duplicate edges (I,J and J,I both appear). If desired, these can be removed by further processing. [Related commands:] none :line :link(Christoff) [(Christoff)] Christoff, "RMAT paper", IEEE (2009). :link(Plimpton) [(Plimpton)] Plimpton, "Large-Scale Graph Algorithm in MapReduce with MPI", to appear in Parallel Computing (2011). "MR-MPI WWW Site" _mws - "MR-MPI Documentation" _md - "OINK Documentation" _od - "OINK Commands" _oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line set command :h3 [Syntax:] set keyword value ... :pre one or more keyword/value pairs may be appended :ulb,l keyword = {verbosity} or {timer} or {memsize} or {outofcore} or {scratch} or {prepend} or {substitute} {verbosity} value = setting for creation of MapReduce objects {timer} value = setting for creation of MapReduce objects

{memsize} value = setting for creation of MapReduce objects {outofcore} value = setting for creation of MapReduce objects {scratch} value = setting for creation of MapReduce objects {prepend} value = string to prepend to file/directory path names {substitute} value = 0 or 1 = how to substitute for "%" in path name :ule [Examples:] set verbosity 2 set verbosity 1 timer 1 memsize 16 set scratch /tmp/mr set prepend /scratch%/data substitute 1 [Description:] This command sets global settings which are used in the creation of MR–MPI objects and the underlying MapReduce objects they wrap. Note that many of these setting names have the same meaning they do in the MR–MPI library themselves, as discussed on "this doc page" [../doc/settings.html](#). The settings for the {verbosity}, {timer}, {memsize}. and {outofcore} keywords are used when the "mr"_mr.html command creates a MapReduce object, unless the "mr"_mr.html command itself overrides these global defaults. "Named commands"_command.html can also create MapReduce objects, either when inputting and outputting data, or when the run() method in the named command class invokes certain methods, like create_mr() or copy_mr(). Each time a new MapReduce object is created, these same global settings for the 4 keywords are applied to it. See "this doc page" [_Section_commands.html](#) for more discussion of the input/output options and these methods. The {scratch} keyword is a directory pathname which all MapReduce objects will use for writing temporary files when they operate in out–of–core mode. Every MapReduce object created by OINK will have its scratch directory set to this value, via the fpath() call described on "this doc page" [../doc/settings.html](#). The {prepend} and {substitute} keywords affect how file and directory names are interpreted by OINK. File and directory names are used as input and output options to "named commands"_command.html via the "–i" and "–o" arguments in an input script. Before these path names are passed to the MR–MPI library, e.g. as part of a "map()" [../doc/map.html](#) method, they can have a directory name prepended to them, and "%" characters in the path name substituted for with a processor ID. This is to enable flexible options for input/output of different files by different processors. If the {prepend} keyword is set, its value should be a directory name (without the trailing "/"). This will be prepended to every input and output pathname used by OINK, including the scratch directory noted above. This global setting can be overridden for a single input or output of the next–executed "named command"_command.html by setting the same {prepend} keyword in the "input"_input.html or "output"_output.html command. Input file or directory names can contain the wildcard character "%". Only the first occurrence of the wildcard character is replaced. If the {substitute} keyword is set to 0, then a "%" is replaced by the processor ID, 0 to Nprocs–1. If it is set to N > 0, then "%" is replaced by (proc–ID % N) + 1. I.e. for 8 processors and N = 4, then the 8 processors replace the "%" with (1,2,3,4,1,2,3,4). This can be useful for multi–core nodes where each core has its own local disk. E.g. you wish each core to read data from one disk. As with the {prepend} keyword, this substitution rule will be applied to every input and output pathname used by OINK, including the scratch directory noted above. This global setting can be overridden for a single input or output of the next–executed "named command"_command.html by setting the same {substitute} keyword in the "input"_input.html or "output"_output.html command. [Related commands:] "input"_input.html, "output"_output.html, "named commands"_command.html, "MR–MPI library commands"_mrmpi.html, "Section_commands"_Section_commands.html [Defaults:] The setting defaults are verbosity = 0, timer = 0, memsize = 64, outofcore = 0, scratch = ".", prepend = NULL, and substitute = 0. "MR–MPI WWW Site"_mws – "MR–MPI Documentation"_md – "OINK Documentation"_od – "OINK Commands"_oc :c :link(mws,<http://www.sandia.gov/~sjplimp/mapreduce.html>) :link(md,[../doc/Manual.html](#)) :link(od,[Manual.html](#)) :link(oc,[Section_commands.html#comm](#)) :line sgi_enumerate command :h3 [Syntax:] sgi_enumerate N V0 F0 E01 V1 F1 E12 V2 F2 ... En–2n–1 Vn–1 Fn–1 –i input1 input2 –o out1.file out1.mr :pre N = number of vertices in sub–graph tour Vi = attribute of each vertex in tour Fi = 0 to N–1 if Vi is same as earlier Vf, –1 if Vi is different than all earlier vertices Eij = attribute of each edge in tour input1 = graph vertices: Key = Vi, Value = label input1 = graph edges: Key = Vi Vj, Value = label out1 = sub–graph matches: Key = Vi Vj ... Vn, Value = NULL :ul [Examples:] sgi_enumerate 4 1 –1 3 2 –1 3 1 –1 3 1 0 –i verts edges –o sgi.list mrsgi :pre [Description:] This is a named command which finds all the sub–graph matches in a large undirected labeled graph. Labeled means that each vertex and edge has an integer value assigned to it. See the "named command"_command.html doc page for various ways in which the –i inputs and –o outputs for a named command can be specified. The labeled sub–graph to search for is specified by the N, Vi, Fi, Eij arguments. Think of the small graph as a tour of edges, walked one by one that includes all the vertices and edges of the sub–graph. Some may need to be listed more than once to describe the graph as a linear walk. The number of vertices in the tour is {N}. The label of each vertex and edge is listed in the tour as {Vi} and {Eij}. {Fi} is used to

flag a vertex if it is encountered again in the tour, meaning that the linear walk has come back on itself. In the example above, triangles are being searched for. This is a list of 4 vertices where the last vertex is the same as the start vertex, thus $F_3 = 0$. All 3 edges in the triangle have a label of 3, so $E_{ij} = 3$. Two of the vertices have a label of 1 (so $V_0 = V_2 = V_3 = 1$), the other has a label of 2 (so $V_1 = 2$). Input1 stores the labeled vertices of the large graph. Input1 stores the labeled edges of the large graph, assumed to have no duplicates or self-edges. The two inputs are unchanged by this command. A list of all matching sub-graphs found in the large graph is calculated. Note that depending on the two graphs, this number can grow very large. Out1 will store the list of sub-graph matches, each as a list of N vertices, which correspond to the {N} vertices in the specified tour. [Related commands:] "sgi_prune"_sgi_prune.html, "sgi_sample"_sgi_sample.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c

:link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line sgi_prune command :h3 [Syntax:] sgi_prune -i input1 input2 -o out1.file out1.mr :pre input1 = graph edges: Key = Vi Vj, Value = NULL input2 = sub-graphs: Key = Vi Vj ... Vn, Value = NULL out1 = new graph edges: Key = Vi Vj, Value = NULL :ul [Examples:] sgi_prune -i mre mrsgi -o graph.list NULL [Description:] This is a named command which uses a set of sub-graphs to remove edges from an undirected graph. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of edges, assumed to have no duplicates or self-edges. Input2 stores a set of sub-graphs, each listed as a series of vertices which imply edges between adjacent vertices. Thus the subgraph (1,10,23,12) implies edges (1,10), (10,23), (23,12). There may be many edges stored multiple times in input2. The two inputs are unchanged by this command. All edges in input1 that do not appear in input2 are removed from the input1 graph. Out1 will store the new graph, containing only edges in input2. The new graph will have no duplicates. [Related commands:]

"sgi_enumerate"_sgi_enumerate.html, "sgi_sample"_sgi_sample.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c

:link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line sgi_sample command :h3 [Syntax:] sgi_sample N V0 F0 E01 V1 F1 E12 V2 F2 ... En-2n-1 Vn-1 Fn-1 -i input1 input2 :pre N = number of vertices in sub-graph tour Vi = attribute of each vertex in tour Fi = 0 to N-1 if Vi is same as earlier Vf, -1 if Vi is different than all earlier vertices Eij = attribute of each edge in tour input1 = graph vertices: Key = Vi, Value = label input1 = graph edges: Key = Vi Vj, Value = label :ul [Examples:] sgi_sample 4 1 -1 3 2 -1 3 1 -1 3 1 0 -i verts edges :pre [Description:] This is a named command which counts all the sub-graph matches in a large undirected labeled graph. Labeled means that each vertex and edge has an integer value assigned to it. Eventually this command will also offer an option to sample and find a subset of the (large) number of matches. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. The labeled sub-graph to search for is specified by the N, Vi, Fi, Eij arguments. Think of the small graph as a tour of edges, walked one by one that includes all the vertices and edges of the sub-graph. Some may need to be listed more than once to describe the graph as a linear walk. The number of vertices in the tour is {N}. The label of each vertex and edge is listed in the tour as {Vi} and {Eij}. {Fi} is used to flag a vertex if it is encountered again in the tour, meaning that the linear walk has come back on itself. In the example above, triangles are being searched for. This is a list of 4 vertices where the last vertex is the same as the start vertex, thus $F_3 = 0$. All 3 edges in the triangle have a label of 3, so $E_{ij} = 3$. Two of the vertices have a label of 1 (so $V_0 = V_2 = V_3 = 1$), the other has a label of 2 (so $V_1 = 2$). Input1 stores the labeled vertices of the large graph. Input1 stores the labeled edges of the large graph, assumed to have no duplicates or self-edges. The two inputs are unchanged by this command. A count of all matching sub-graphs found in the large graph is calculated in an efficient manner that does not require enumerating all of them, like the "sgi_enumerate"_sgi_enumerate.html command does. Note that depending on the two graphs, this number can grow very large. This command currently has no outputs. Eventually it will allow a sampled subset of the sub-graph matches to be output. [Related commands:]

"sgi_prune"_sgi_prune.html, "sgi_enumerate"_sgi_enumerate.html "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c

:link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line shell command :h3 [Syntax:] shell style args :pre style = {cd} or {mkdir} or {mv} or {rm} or {rmdir} :ulb,1 {cd} arg = dir dir = directory to change to {mkdir} args = dir1 dir2 ...

dir1,dir2 = one or more directories to create {mv} args = old new old = old filename new = new filename {rm} args = file1 file2 ... file1,file2 = one or more filenames to delete {rmdir} args = dir1 dir2 ... dir1,dir2 = one or more directories to delete :pre :ule [Examples:] shell cd sub1 shell cd .. shell mkdir tmp1 tmp2 tmp3 shell rmdir tmp1 shell mv log.lammps hold/log.1 shell rm TMP/file1 TMP/file2 :pre [Description:] Execute a shell command. Only a few simple file-based shell commands are supported, in Unix-style syntax. With the exception of {cd}, all commands are executed by only a single processor, so that files/directories are not being manipulated by multiple processors. The {cd} style executes the Unix "cd" command to change the working directory. All subsequent OINK commands that read/write files will use the new directory. All processors execute this command. The {mkdir} style executes the Unix "mkdir" command to create one or more directories. The {mv} style executes the Unix "mv" command to rename a file and/or move it to a new directory. The {rm} style executes the Unix "rm" command to remove one or more files. The {rmdir} style executes the Unix "rmdir" command to remove one or more directories. A directory must be empty to be successfully removed.

IMPORTANT NOTE: OINK does not detect errors or print warnings when any of these Unix commands execute. E.g. if the specified directory does not exist, executing the {cd} command will silently not do anything. [Related commands:] none "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line tri_find command :h3 [Syntax:] tri_find -i input1 -o out1.file out1.mr :pre input1 = graph edges: Key = Vi Vj, Value = NULL out1 = triangles: Key = Vi Vj Vk, Value = NULL :ul [Examples:] tri_find -i mre -o tri.list mtri [Description:] This is a named command which finds all the triangles in an undirected graph. A triangle is a set of 3 vertices I,J,K for which the edges IJ, JK, IK all exist. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of edges, assumed to have no duplicates or self-edges. The input is unchanged by this command. The triangles are found via the MapReduce algorithm of "(Cohen)"_#Cohen discussed in his paper and in the paper of "(Plimpton)"_#Plimpton. Note that even small graphs can have large numbers of triangles if there are very high-degree vertices. Out1 will store the list of triangles. [Related commands:] none :line :link(Cohen) [(Cohen)] Cohen, "Graph Twiddling in a MapReduce World", IEEE (2009). :link(Plimpton) [(Plimpton)] Plimpton, "Large-Scale Graph Algorithm in MapReduce with MPI", to appear in Parallel Computing (2011). "MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,..doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line variable command :h3 [Syntax:] variable name style args ... :pre name = name of variable to define :ulb,l style = {delete} or {index} or {loop} or {world} or {universe} or {uloop} or {string} or {equal} or {atom} :l {delete} = no args {index} args = one or more strings {loop} args = N N = integer size of loop, loop from 1 to N inclusive {loop} args = N pad N = integer size of loop, loop from 1 to N inclusive pad = all values will be same length, e.g. 001, 002, ..., 100 {loop} args = N1 N2 N1,N2 = loop from N1 to N2 inclusive {loop} args = N1 N2 pad N1,N2 = loop from N1 to N2 inclusive pad = all values will be same length, e.g. 050, 051, ..., 100 {world} args = one string for each partition of processors {universe} args = one or more strings {uloop} args = N N = integer size of loop {uloop} args = N pad N = integer size of loop pad = all values will be same length, e.g. 001, 002, ..., 100 {string} arg = one string {equal} args = one formula containing numbers, keywords, math operations, variable references numbers = 0.0, 100, -5.4, 2.8e-4, etc constants = PI keywords = nprocs, time math operators = (), -x, x+y, x-y, x*y, x/y, x^y, x==y, x!=y, xy, x>=y, xx||y, !x math functions = sqrt(x), exp(x), ln(x), log(x), sin(x), cos(x), tan(x), asin(x), acos(x), atan(x), atan2(y,x), random(x,y,z), normal(x,y,z), ceil(x), floor(x), round(x) ramp(x,y), stagger(x,y), logfreq(x,y,z), vdisplace(x,y), swiggle(x,y,z), cwiggle(x,y,z) variable references = v_name :pre :ule [Examples:] variable x index run1 run2 run3 run4 run5 run6 run7 run8 variable LoopVar loop \$n variable p equal nprocs variable b1 equal 0.5*v_flag variable b1 equal "10 + 0.5*v_flag" variable foo myfile variable x universe 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 variable x uloop 15 pad variable x delete :pre [Description:] This command assigns one or more strings to a variable name for evaluation later in the input script or during a simulation. Variables can be used in several ways in OINK. A variable can be referenced elsewhere in an input script to become part of a new input command. For variable styles that store multiple strings, the "next"_next.html command can be used to increment which string is assigned to the variable. Variables of style {equal} store a formula which when evaluated produces a single numeric value which can be output via the "print"_print.html command. Variables that store a collection of strings can be used as input to a

named command, e.g. to process a collection of filenames. See the "named command"_command.html doc page for details. In the discussion that follows, the "name" of the variable is the arbitrary string that is the 1st argument in the variable command. This name can only contain alphanumeric characters and underscores. The "string" is one or more of the subsequent arguments. The "string" can be simple text as in the 1st example above, it can contain other variables as in the 2nd example, or it can be a formula as in the 3rd example. The "value" is the numeric quantity resulting from evaluation of the string. Note that the same string can generate different values when it is evaluated at different times during a simulation. **IMPORTANT NOTE:** When the input script line that defines a variable of style {equal} is encountered, the formula is NOT immediately evaluated and the result stored. See the discussion below about "Immediate Evaluation of Variables" if you want to do this. **IMPORTANT NOTE:** When a variable command is encountered in the input script and the variable name has already been specified, the command is ignored. This means variables can NOT be re-defined in an input script (with 2 exceptions, read further). This is to allow an input script to be processed multiple times without resetting the variables; see the "jump"_jump.html or "include"_include.html commands. It also means that using the "command-line switch" _Section_start.html#2_6 -var will override a corresponding index variable setting in the input script. There are two exceptions to this rule. First, variables of style {string} and {equal} ARE redefined each time the command is encountered. This allows these style of variables to be redefined multiple times in an input script. In a loop, this means the formula associated with an {equal}-style variable can change if it contains a substitution for another variable, e.g. \$x. Second, as described below, if a variable is iterated on to the end of its list of strings via the "next"_next.html command, it is removed from the list of active variables, and is thus available to be re-defined in a subsequent variable command. The {delete} style does the same thing. :line "This section" _Section_script.html#2_2 of the manual explains how occurrences of a variable name in an input script line are replaced by the variable's string. The variable name can be referenced as \$x if the name "x" is a single character, or as \${LoopVar} if the name "LoopVar" is one or more characters. As described below, for variable styles {index}, {loop}, {universe}, and {uloop}, which string is assigned to a variable can be incremented via the "next"_next.html command. When there are no more strings to assign, the variable is exhausted and a flag is set that causes the next "jump"_jump.html command encountered in the input script to be skipped. This enables the construction of simple loops in the input script that are iterated over and then exited from. As explained above, an exhausted variable can be re-used in an input script. The {delete} style also removes the variable, the same as if it were exhausted, allowing it to be redefined later in the input script or when the input script is looped over. This can be useful when breaking out of a loop via the "if"_if.html and "jump"_jump.html commands before the variable would become exhausted. For example, label loop variable a loop 5 print "A = \$a" if \$a > 2 then "jump in.script break" next a jump in.script loop label break variable a delete :pre :line For the {index} style, one or more strings are specified. Initially, the 1st string is assigned to the variable. Each time a "next"_next.html command is used with the variable name, the next string is assigned. All processors assign the same string to the variable. {Index} style variables with a single string value can also be set by using the command-line switch -var; see "this section" _Section_start.html#2_6 for details. The {loop} style is identical to the {index} style except that the strings are the integers from 1 to N inclusive, if only one argument N is specified. This allows generation of a long list of runs (e.g. 1000) without having to list N strings in the input script. Initially, the string "1" is assigned to the variable. Each time a "next"_next.html command is used with the variable name, the next string ("2", "3", etc) is assigned. All processors assign the same string to the variable. The {loop} style can also be specified with two arguments N1 and N2. In this case the loop runs from N1 to N2 inclusive, and the string N1 is initially assigned to the variable. For the {world} style, one or more strings are specified. There must be one string for each processor partition or "world". See "this section" _Section_build.html#1_4 of the manual for information on running OINK with multiple partitions via the "-partition" command-line switch. This variable command assigns one string to each world. All processors in the world are assigned the same string. The next command cannot be used with {equal} style variables, since there is only one value per world. This style of variable is useful when you wish to perform different calculations on different partitions. For the {universe} style, one or more strings are specified. There must be at least as many strings as there are processor partitions or "worlds". See "this page" _Section_start.html#2_6 for information on running OINK with multiple partitions via the "-partition" command-line switch. This variable command initially assigns one string to each world. When a "next"_next.html command is encountered using this variable, the first processor partition to encounter it, is assigned the next available string. This continues until all the variable strings are consumed. Thus, this command

can be used to run 50 simulations on 8 processor partitions. The simulations will be run one after the other on whatever partition becomes available, until they are all finished. {Universe} style variables are incremented using the files "tmp.oink.variable" and "tmp.oink.variable.lock" which you will see in your directory during such a OINK run. The {uloop} style is identical to the {universe} style except that the strings are the integers from 1 to N. This allows generation of long list of runs (e.g. 1000) without having to list N strings in the input script. All {universe}- and {uloop}-style variables defined in an input script must have the same number of values. :line

For the {equal} style, a single string is specified which represents a formula that will be evaluated afresh each time the variable is used. If you want spaces in the string, enclose it in double quotes so the parser will treat it as a single argument. The formula computes a scalar quantity, which becomes the value of the variable whenever it is evaluated. Note that {equal} variables can produce different values at different stages of the input script or at different times during a run. For example, if the {equal} variable is printed during a loop, different values could be printed each time it was invoked. If you want a variable to be evaluated immediately, so that the result is stored by the variable instead of the string, see the section below on "Immediate Evaluation of Variables". The next command cannot be used with {equal} style variables, since there is only one string. The formula for an {equal} variable can contain a variety of quantities. The syntax for each kind of quantity is simple, but multiple quantities can be nested and combined in various ways to build up formulas of arbitrary complexity. Specifically, an formula can contain numbers, keywords, math operators, math functions, and references to other variables.

Number: 0.2, 100, 1.0e20, -15.4, etc Constant: PI Keywords: nprocs, time Math operators: (), -x, x+y, x-y, x*y, x/y, x^y, x==y, x!=y, xy, x>=y, xx||y, !x Math functions: sqrt(x), exp(x), ln(x), log(x), sin(x), cos(x), tan(x), asin(x), acos(x), atan(x), atan2(y,x), random(x,y,z), normal(x,y,z), ceil(x), floor(x), round(x), ramp(x,y), stagger(x,y), logfreq(x,y,z), vdisplace(x,y), swiggle(x,y,z), cwiggle(x,y,z) Other variables: v_name :tb(s=) :line

The keywords allowed in a formula are {nprocs} and {time}. Nprocs is the number of processors being used. Time is the elapsed time of the most recently executed "named command"_command.html or "MR-MPI library command"_mrmpi.html. :line

Math Operators :h4 Math operators are written in the usual way, where the "x" and "y" in the examples can themselves be arbitrarily complex formulas, as in the examples above. Operators are evaluated left to right and have the usual C-style precedence: unary minus and unary logical NOT operator "!" have the highest precedence, exponentiation "^" is next; multiplication and division are next; addition and subtraction are next; the 4 relational operators "=", ">=", "<=", and "<=" are next; the two remaining relational operators "==" and "!=" are next; then the logical AND operator "&" and finally the logical OR operator "||" has the lowest precedence. Parenthesis can be used to group one or more portions of a formula and/or enforce a different order of evaluation than what would occur with the default precedence. The 6 relational operators return either a 1.0 or 0.0 depending on whether the relationship between x and y is TRUE or FALSE. For example the expression

x"MR-MPI WWW Site"_mws -"MR-MPI Documentation"_md - "OINK Documentation"_od - "OINK Commands"_oc :c :link(mws,http://www.sandia.gov/~sjplimp/mapreduce.html) :link(md,./doc/Manual.html) :link(od,Manual.html) :link(oc,Section_commands.html#comm) :line

wordfreq command :h3 [Syntax:] wordfreq Ntop -i input1 -o out1.file out1.mr :pre Ntop = print Ntop of the most frequently occurring words to screen

input1 = words: Key = word, Value = NULL out1 = frequency count of each word: Key = word, Value = count :ul [Examples:] wordfreq 10 -i v_files -o full.list NULL wordfreq 10 -i v_files -o NULL NULL [Description:]

This is a named command which calculates the frequency of word occurrence in an input data set, which is typically a set of files. See the "named command"_command.html doc page for various ways in which the -i inputs and -o outputs for a named command can be specified. Input1 stores a set of words. The input is changed by the command as it is processed further. If the input is one or more files then the files are read and each "word" is defined as separated by whitespace. Note that you can pass a list of files as the input argument after the "-i" argument by using a variable, which in turn can be initialized with a command-line argument to OINK. E.g. this line would work with the first example above: oink_linux -var files *.cpp < in.script :pre

See "this section"_Section_build.html#1_4 of the manual and the "variable"_variable.html doc page for more details. A count of the number of times each word occurs is calculated. The highest frequency {Ntop} occurrences are printed to the screen with their count. If {Ntop} is 0, nothing is printed. Out1 will store the frequency count of all unique words. [Related commands:] none